
Reading Great Code

Programmers read a *lot* of code. One of the core tenets behind Python’s design is readability, and one secret to becoming a great programmer is to read, understand, and comprehend excellent code. Such code typically follows the guidelines outlined in “Code Style” on page 43 and does its best to express a clear and concise intent to the reader.

This chapter shows excerpts from some very readable Python projects that illustrate topics covered in Chapter 4. As we describe them, we’ll also share techniques for reading code.¹

Here’s a list of projects highlighted in this chapter in the order they will appear:

- **HowDoI** is a console application that searches the Internet for answers to coding questions, written in Python.
- **Diamond** is a Python daemon² that collects metrics and publishes them to Graphite or other backends. It is capable of collecting CPU, memory, network, I/O, load and disk metrics. Additionally, it features an API for implementing custom collectors to gather metrics from almost any source.
- **Tablib** is a format-agnostic tabular dataset library.
- **Requests** is a HyperText Transfer Protocol (HTTP) library for human beings (the 90% of us who just want an HTTP client that automatically handles password

¹ For a book that contains decades of experience about reading and refactoring code, we recommend *Object-Oriented Reengineering Patterns* (Square Bracket Associates) by Serge Demeyer, Stéphane Ducasse, and Oscar Nierstrasz.

² A *daemon* is a computer program that runs as a background process.

authentication and complies with the **half-dozen standards** to perform things like a multipart file upload with one function call).

- **Werkzeug** started as a simple collection of various utilities for Web Service Gateway Interface (WSGI) applications and has become one of the most advanced WSGI utility modules.
- **Flask** is a web microframework for Python based on Werkzeug and Jinja2. It's good for getting simple web pages up quickly.

There is a lot more to all of these projects than what we're mentioning, and we really, really hope that after this chapter you'll be motivated to download and read at least one or two of them in depth yourself (and maybe even present what you learn to a local user group).

Common Features

Some features are common across all of the projects: details from a snapshot of each one show very few (fewer than 20, excluding whitespace and comments) lines of code on average per function, and a lot of blank lines. The larger, more complex projects use docstrings and/or comments; usually more than a fifth of the content of the code base is some sort of documentation. But we can see from HowDoI, which has no docstrings because it is not for interactive use, that comments are not necessary when the code is straightforward. **Table 5-1** shows common practices in these projects.

Table 5-1. Common features in the example projects

Package	License	Line count	Docstrings (% of lines)	Comments (% of lines)	Blank lines (% of lines)	Average function length
HowDoI	MIT	262	0%	6%	20%	13 lines of code
Diamond	MIT	6,021	21%	9%	16%	11 lines of code
Tablib	MIT	1,802	19%	4%	27%	8 lines of code
Requests	Apache 2.0	4,072	23%	8%	19%	10 lines of code
Flask	BSD 3-clause	10,163	7%	12%	11%	13 lines of code
Werkzeug	BSD 3-clause	25,822	25%	3%	13%	9 lines of code

In each section, we use a different code-reading technique to figure out what the project is about. Next, we single out code excerpts that demonstrate ideas mentioned elsewhere in this guide. (Just because we don't highlight things in one project doesn't mean they don't exist; we just want to provide good coverage of concepts across these examples.) You should finish this chapter more confident about reading code, with examples that reinforce what makes good code, and with some ideas you'd like to incorporate in your own code later.

HowDoI

With fewer than 300 lines of code, The HowDoI project, by Benjamin Gleitzman, is a great choice to start our reading odyssey.

Reading a Single-File Script

A script usually has a clear starting point, clear options, and a clear ending point. This makes it easier to follow than libraries that present an API or provide a framework.

Get the HowDoI module from GitHub:³

```
$ git clone https://github.com/gleitz/howdoi.git
$ virtualenv -p python3 venv # or use mkvirtualenv, your choice...
$ source venv/bin/activate
(venv)$ cd howdoi/
(venv)$ pip install --editable .
(venv)$ python test_howdoi.py # Run the unit tests.
```

You should now have the howdoi executable installed in *venv/bin*. (You can look at it if you want by typing `cat `which howdoi`` on the command line.) It was auto-generated when you ran `pip install`.

Read HowDoI's documentation

HowDoI's documentation is in the *README.rst* file in the [HowDoI repository on GitHub](#): it's a small command-line application that allows users to search the Internet for answers to programming questions.

From the command line in a terminal shell, we can type `howdoi --help` for the usage statement:

```
(venv)$ howdoi --help
usage: howdoi [-h] [-p POS] [-a] [-l] [-c] [-n NUM_ANSWERS] [-C] [-v]
              [QUERY [QUERY ...]]
```

instant coding answers via the [command line](#)

positional arguments:

QUERY the question to answer

optional arguments:

-h, --help	show this help message and exit
-p POS, --pos POS	select answer in specified position (default: 1)
-a, --all	display the full text of the answer
-l, --link	display only the answer link

³ If you run into trouble with `lxml` requiring a more recent `libxml2` shared library, just install an earlier version of `lxml` by typing: `pip uninstall lxml; pip install lxml==3.5.0`. It will work fine.

```

-c, --color          enable colorized output
-n NUM_ANSWERS, --num-answers NUM_ANSWERS
                    number of answers to return
-C, --clear-cache    clear the cache
-v, --version        displays the current version of howdoi

```

That's it—from the documentation we know that HowDoI gets answers to coding questions from the Internet, and from the usage statement we know we can choose the answer in a specific position, can colorize the output, get multiple answers, and that it keeps a cache that can be cleared.

Use HowDoI

We can confirm we understand what HowDoI does by actually using it. Here's an example:

```

(venv)$ howdoi --num-answers 2 python lambda function list comprehension
--- Answer 1 ---
[(lambda x: x*x)(x) for x in range(10)]

--- Answer 2 ---
[x() for x in [lambda m=m: m for m in [1,2,3]]]
# [1, 2, 3]

```

We've installed HowDoI, read its documentation, and can use it. On to reading actual code!

Read HowDoI's code

If you look inside the *howdoi/* directory, you'll see it contains two files: an *__init__.py*, which contains a single line that defines the version number, and *howdoi.py*, which we'll open and read.

Skimming *howdoi.py*, we see each new function definition is used in the next function, making it is easy to follow. And each function does just one thing—the thing its name says. The main function, *command_line_runner()*, is near the bottom of *howdoi.py*.

Rather than reprint HowDoI's source here, we can illustrate its call structure using the call graph in [Figure 5-1](#). It was created by [Python Call Graph](#), which provides a visualization of the functions called when running a Python script. This works well with command-line applications thanks to a single start point and the relatively few paths through their code. (Note that we manually deleted functions not in the HowDoI project from the rendered image to legibly fit it on the page, and slightly recolored and reformatted it.)

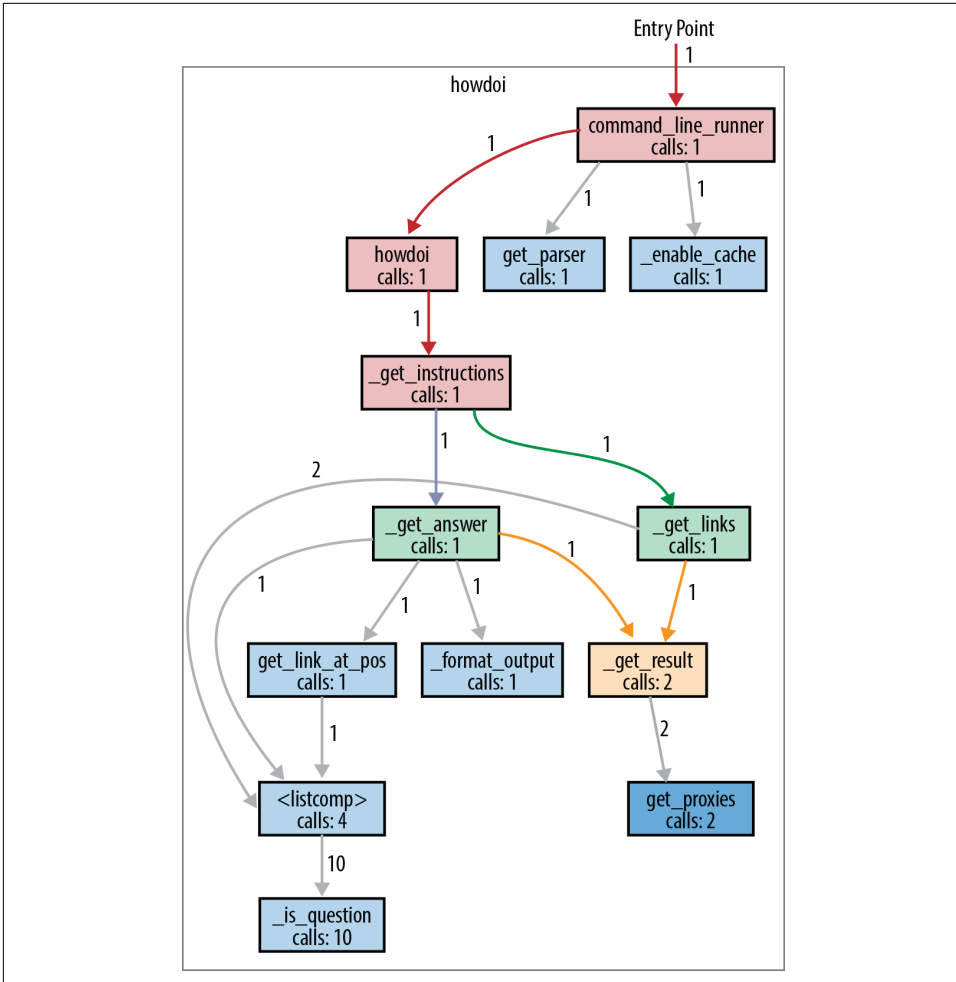


Figure 5-1. Clean paths and clear function names in this howdoi call graph

The code could have been all one large, incomprehensible spaghetti function. Instead, intentional choices structure the code into compartmentalized functions with straightforward names. Here's a brief description of the execution depicted in [Figure 5-1](#): `command_line_runner()` parses the input and passes the user flags and the query to `howdoi()`. Then, `howdoi()` wraps `_get_instructions()` in a try/except statement so that it can catch connection errors and print a reasonable error message (because application code should not terminate on exceptions).

The primary functionality is in `_get_instructions()`: it calls `_get_links()` to do a Google search of Stack Overflow for links that match the query, then calls

`_get_answer()` once for each resulting link (up to the number of links that the user specified on the command line—the default is just one link).

The `_get_answer()` function follows a link to Stack Overflow, extracts code from the answer, colorizes it, and returns it to `_get_instructions()`, which will combine all of the answers into one string, and return it. Both `_get_links()` and `_get_answer()` call `_get_result()` to actually do the HTTP request: `_get_links()` for the Google query, and `_get_answer()` for the resulting links from the Google query.

All `_get_result()` does is wrap `requests.get()` with a `try/except` statement so that it can catch SSL errors, print an error message, and re-raise the exception so that the top-level `try/except` can catch it and exit. Catching all exceptions before exiting is best practice for application programs.

HowDoI's Packaging

HowDoI's `setup.py`, above the `howdoi/` directory, is a good example setup module because in addition to normal package installation, it also installs an executable (which you can refer to when packaging your own command-line utility). The `setup.tools.setup()` function uses keyword arguments to define all of the configuration options. The part that identifies the executable is associated with the keyword argument `entry_points`:

```
setup(
    name='howdoi',
    ##~~ ... Skip the other typical entries ...
    entry_points={
        'console_scripts': [ ❶
            'howdoi = howdoi.howdoi:command_line_runner', ❷
        ],
    },
    ## ~~ ... Skip the list of dependencies ...
)
```

- ❶ The keyword to list console scripts is `console_scripts`.
- ❷ This declares the executable named `howdoi` will have as its target the function `howdoi.howdoi.command_line_runner()`. So later when reading, we will know `command_line_runner()` is the starting point for running the whole application.

Structure Examples from HowDoI

HowDoI is a small library, and we'll be highlighting structure much more elsewhere, so there are only a few notes here.

Let each function do just one thing

We can't reiterate enough how beneficial it is for readers to separate out HowDoI's internal functions to each do just one thing. Also, there are functions whose sole purpose is to wrap other functions with a `try/except` statement. (The only function with a `try/except` that doesn't follow this practice is `_format_output()`, which leverages `try/except` clauses to identify the correct coding language for syntax highlighting, not for exception handling.)

Leverage data available from the system

HowDoI checks and uses relevant system values, such as `urllib.request.getproxies()`, to handle the use of proxy servers (this can be the case in organizations like schools that have an intermediary server filtering the connection to the Internet), or in this snippet:

```
XDG_CACHE_DIR = os.environ.get(
    'XDG_CACHE_HOME',
    os.path.join(os.path.expanduser('~'), '.cache')
)
```

How do you know that these variables exist? The need for `urllib.request.getproxies()` is evident from the optional arguments in `requests.get()`—so part of this information comes from understanding the API of libraries you call. Environment variables are often utility-specific, so if a library is intended for use with a particular database or other sister application, those applications' documentation list relevant environment variables. For plain POSIX systems, a good place to start is [Ubuntu's list of default environment variables](#), or else the base list of environment variables in the [POSIX specification](#), which links to various relevant other lists.

Style Examples from HowDoI

HowDoI mostly follows PEP 8, but not pedantically, and not when it restricts readability. For example, `import` statements are at the top of the file, but standard library and external modules are intermixed. And although the string constants in `USER_AGENTS` are much longer than 80 characters, there is no natural place to break the strings, so they are left intact.

These next excerpts highlight other style choices we've previously advocated for in [Chapter 4](#).

Underscore-prefixed function names (we are all responsible users)

Almost every function in HowDoI is prefixed with an underscore. This identifies them as for internal use only. For most of them, this is because if called, there is the

possibility of an uncaught exception—anything that calls `_get_result()` risks this—until the `howdoi()` function, which handles the possible exceptions.

The rest of the internal functions (`_format_output()`, `_is_question()`, `_enable_cache()`, and `_clear_cache()`) are identified as such because they’re simply not intended for use outside of the package. The testing script, *howdoi/test_howdoi.py*, only calls the nonprefixed functions, checking that the formatter works by feeding a command-line argument for colorization to the top-level `howdoi.howdoi()` function, rather than by feeding code to `howdoi._format_output()`.

Handle compatibility in just one place (readability counts)

Differences between versions of possible dependencies are handled before the main code body so the reader knows there won’t be dependency issues, and version checking doesn’t litter the code elsewhere. This is nice because HowDoI is shipped as a command-line tool, and the extra effort means users won’t be forced to change their Python environment just to accommodate the tool. Here is the snippet with the workarounds:

```
try:
    from urllib.parse import quote as url_quote
except ImportError:
    from urllib import quote as url_quote

try:
    from urllib import getproxies
except ImportError:
    from urllib.request import getproxies
```

And the following snippet resolves the difference between Python 2 and Python 3’s Unicode handling in seven lines, by creating the function `u(x)` to either do nothing or emulate Python 3. Plus it follows Stack Overflow’s [new citation guideline](#), by citing the original source:

```
# Handle Unicode between Python 2 and 3
# http://stackoverflow.com/a/6633040/305414
if sys.version < '3':
    import codecs
    def u(x):
        return codecs.unicode_escape_decode(x)[0]
else:
    def u(x):
        return x
```

Pythonic choices (beautiful is better than ugly)

The following snippet from *howdoi.py* shows thoughtful, Pythonic choices. The function `get_link_at_pos()` returns `False` if there are no results, or else identifies the

links that are to Stack Overflow questions, and returns the one at the desired position (or the last one if there aren't enough links):

```
def _is_question(link): ❶
    return re.search('questions/\d+/', link)

# [ ... skip a function ... ]

def get_link_at_pos(links, position):
    links = [link for link in links if _is_question(link)] ❷
    if not links:
        return False ❸

    if len(links) >= position:
        link = links[position-1] ❹
    else:
        link = links[-1] ❺
    return link ❻
```

- ❶ The first function, `_is_question()`, is defined as a separate one liner, giving clear meaning to an otherwise opaque regular expression search.
- ❷ The list comprehension reads like a sentence, thanks to the separate definition of `_is_question()` and meaningful variable names.
- ❸ The early `return` statement flattens the code.
- ❹ The additional step of assigning to the variable `link` here...
- ❺ ...and here, rather than two separate `return` statements with no named variable at all, reinforces the purpose of `get_link_at_pos()` with clear variable names. The code is self-documenting.
- ❻ The single `return` statement at the highest indentation level explicitly shows that all paths through the code exit either right away—because there are no links—or at the end of the function, returning a link. Our quick rule of thumb works: we can read the first and last line of this function and understand what it does. (Given multiple links and a position, `get_link_at_pos()` returns one single link: the one at the given position.)

Diamond

Diamond is a daemon (an application that runs continuously as a background process) that collects system metrics and publishes them to downstream programs like MySQL, **Graphite** (a platform open sourced by Orbitz in 2008 that stores, retrieves,

and optionally graphs numeric time-series data), and others. We'll get to explore good package structure, as Diamond is a multifile application, much larger than HowDoI.

Reading a Larger Application

Diamond is still a command-line application, so like with HowDoI, there's still a clear starting point and clear paths of execution, although the supporting code now spans multiple files.

Get Diamond from GitHub (the documentation says it only runs on CentOS or Ubuntu, but code in its *setup.py* makes it appear to support all platforms; however, some of the commands that default collectors use to monitor memory, disk space, and other system metrics are not on Windows). As of this writing, it still uses Python 2.7:

```
$ git clone https://github.com/python-diamond/Diamond.git
$ virtualenv -p python2 venv # It's not Python 3 compatible yet...
$ source venv/bin/activate
(venv)$ cd Diamond/
(venv)$ pip install --editable .
(venv)$ pip install mock docker-py # These are dependencies for testing.
(venv)$ pip install mock # This is also a dependency for testing.
(venv)$ python test.py # Run the unit tests.
```

Like with the HowDoI library, Diamond's setup script installs executables in *venv/bin/*: *diamond* and *diamond-setup*. This time they're not automatically generated—they're prewritten scripts in the project's *Diamond/bin/* directory. The documentation says that *diamond* starts the server, and *diamond-setup* is an optional tool to walk users through interactive modification of the collector settings in the configuration file.

There are a lot of additional directories, and the *diamond* package is underneath *Diamond/src* in this project directory. We are going to look at files in *Diamond/src* (which contains the main code), *Diamond/bin* (which contains the executable *diamond*), and *Diamond/conf* (which contains the sample configuration file). The rest of the directories and files may be of interest to people distributing similar applications but is not what we want to cover right now.

Read Diamond's documentation

First, we can get a sense of what the project is and what it does by scanning the [online documentation](#). Diamond's goal is to make it easy to gather system metrics on clusters of machines. Originally open sourced by BrightCove, Inc., in 2011, it now has over 200 contributors.

After describing its history and purpose, the documentation tells you how to install it, and then says how to run it: just modify the example configuration file (in our down-

load it's in *conf/diamond.conf.example*), put it in the default location (*/etc/diamond/diamond.conf*) or a path you'll specify on the command line, and you're set. There's also a helpful section on configuration in the [Diamond wiki page](#).

From the command line, we can get the usage statement via `diamond --help`:

```
(venv)$ diamond --help
Usage: diamond [options]

Options:
  -h, --help            show this help message and exit
  -c CONFIGFILE, --configfile=CONFIGFILE
                        config file
  -f, --foreground      run in foreground
  -l, --log-stdout      log to stdout
  -p PIDFILE, --pidfile=PIDFILE
                        pid file
  -r COLLECTOR, --run=COLLECTOR
                        run a given collector once and exit
  -v, --version          display the version and exit
  --skip-pidfile        Skip creating PID file
  -u USER, --user=USER  Change to specified unprivileged user
  -g GROUP, --group=GROUP
                        Change to specified unprivileged group
  --skip-change-user    Skip changing to an unprivileged user
  --skip-fork           Skip forking (daemonizing) process
```

From this, we know it uses a configuration file; by default, it runs in the background; it has logging; you can specify a PID (process ID) file; you can test collectors; you can change the process's user and group; and it by default will daemonize (fork) the process.⁴

Use Diamond

To understand it even better, we can run Diamond. We need a modified configuration file, which we can put in a directory we make called *Diamond/tmp*. From inside the *Diamond* directory, type:

```
(venv)$ mkdir tmp
(venv)$ cp conf/diamond.conf.example tmp/diamond.conf
```

Then edit *tmp/diamond.conf* to look like this:

⁴ When you daemonize a process, you fork it, detach its session ID, and fork it again, so that the process is totally disconnected from the terminal you're running it in. (Nondaemonized programs exit when the terminal is closed—you may have seen the warning message “Are you sure you want to close this terminal? Closing it will kill the following processes:” before listing all of the currently running processes.) A daemonized process will run even after the terminal window closes. It's named daemon after [Maxwell's daemon](#) (a clever daemon, not a nefarious one).

```

### Options for the server
[server]
# Handlers for published metrics. ❶
handlers = diamond.handler.archive.ArchiveHandler
user = ❷
group =
# Directory to load collector modules from ❸
collectors_path = src/collectors/

### Options for handlers ❹
[handlers]
[[default]]

[[ArchiveHandler]]
log_file = /dev/stdout

### Options for collectors
[collectors]
[[default]]
# Default Poll Interval (seconds)
interval = 20

### Default enabled collectors
[[CPUCollector]]
enabled = True

[[MemoryCollector]]
enabled = True

```

We can tell from the example configuration file that:

- ❶ There are multiple handlers, which we can select by class name.
- ❷ We have control over the user and group that the daemon runs as (empty means to use the current user and group).
- ❸ We can specify a path to look for collector modules. This is how Diamond will know where the custom Collector subclasses are: we directly state it in the configuration file.
- ❹ We can also store configure handlers individually.

Next, run Diamond with options that set logging to `/dev/stdout` (with default formatting configurations), that keep the application in the foreground, that skip writing the PID file, and that use our new configuration file:

```
(venv)$ diamond -l -f --skip-pidfile --configfile=tmp/diamond.conf
```

To end the process, type Ctrl+C until the command prompt reappears. The log output demonstrates what collectors and handlers do: collectors collect different metrics

(such as the `MemoryCollector`’s total, available, free, and swap memory sizes), which the handlers format and send to various destinations, such as Graphite, MySQL, or in our test case, as log messages to `/dev/stdout`.

Reading Diamond’s code

IDEs can be useful when reading larger projects—they can quickly locate the original definitions of functions and classes in the source code. Or, given a definition, they can find all places in the project where it is used. For this functionality, set the IDE’s Python interpreter to the one in your virtual environment.⁵

Instead of following each function as we did with HowDoI, **Figure 5-2** follows the `import` statements; the diagram just shows which modules in Diamond import which other modules. Drawing sketches like these helps by providing a very high-level look for larger projects: you hide the trees so you can see the forest. We can start with the `diamond` executable file on the top left and follow the imports through the Diamond project. Aside from the `diamond` executable, every square outline denotes a file (module) or directory (package) in the `src/diamond` directory.

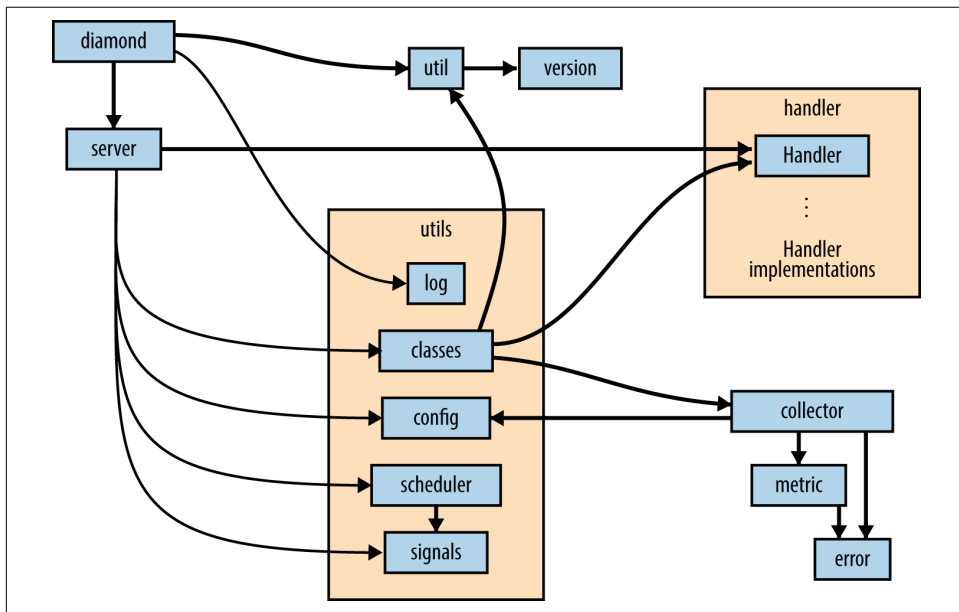


Figure 5-2. The module import structure of Diamond

⁵ In PyCharm, do this by navigating in the menu bar to PyCharm → Preferences → Project:Diamond → Project Interpreter, and then selecting the path to the Python interpreter in the current virtual environment.

Diamond's well-organized and appropriately named modules make it possible to get an idea of what the code is doing solely from our diagram: `diamond` gets the version from `util`, then sets up logging using `utils.log` and starts a `Server` instance using `server`. The `Server` imports from almost all of the modules in the `utils` package, using `utils.classes` to access both the `Handlers` in `handler` and the `collectors`, `config` to read the configuration file and obtain settings for the `collectors` (and the extra paths to the user-defined `collectors`), and `scheduler` and `signals` to set the polling interval for the `collectors` to calculate their metrics, and to set up and start the `handlers` processing the queue of metrics to send them to their various destinations.

The diagram doesn't include the helper modules `convertor.py` and `gmetric.py`, which are used by specific `collectors`, or the over 20 `handler` implementations defined in the `handler` subpackage, or the over 100 `collector` implementations defined in the project's `Diamond/src/collectors/` directory (which is installed elsewhere when not installed the way we did for reading—that is, using PyPI or Linux package distributions, instead of source). These are imported using `diamond.classes.load_dynamic_class()`, which then calls the function `diamond.util.load_class_from_name()` to load the classes from the string names given in the configuration file, so the `import` statements do not explicitly name them.

To understand why there is both a `utils` package and a `util` module, you have to dig into the actual code: the `util` module provides functions related more to Diamond's packaging than to its operation—a function to get the version number from `version.__VERSION__`, and two functions that parse strings that identify either modules or classes, and import them.

Logging in Diamond

The function `diamond.utils.log.setup_logging()`, found in `src/diamond/utils/log.py`, is called from the `main()` function in the `diamond` executable when starting the daemon:

```
# Initialize logging
log = setup_logging(options.configfile, options.log_stdout)
```

If `options.log_stdout` is `True`, `setup_logging()` will set up a logger with default formatting to log to standard output at the `DEBUG` level. Here's the excerpt that does that:

```
##~~ ... Skip everything else ...

def setup_logging(configfile, stdout=False):
    log = logging.getLogger('diamond')

    if stdout:
        log.setLevel(logging.DEBUG)
```

```

streamHandler = logging.StreamHandler(sys.stdout)
streamHandler.setFormatter(DebugFormatter())
streamHandler.setLevel(logging.DEBUG)
log.addHandler(streamHandler)
else:
    ##~~ ... Skip this ...

```

Otherwise, it parses the configuration file using `logging.config.fileConfig()` from the Python Standard Library. Here is the function call—it’s indented because it’s inside the preceding `if/else` statement, and a `try/except` block:

```

logging.config.fileConfig(configfile,
                          disable_existing_loggers=False)

```

The logging configuration ignores keywords in the configuration file that aren’t related to logging. This is how Diamond can use the same configuration file for both its own and the logging configuration. The sample configuration file, located in *Diamond/conf/diamond.conf.example*, identifies the logging handler among the other Diamond handlers:

```

### Options for handlers
[handlers]

# daemon logging handler(s)
keys = rotated_file

```

It defines example loggers later in the configuration file, under the header “Options for logging,” recommending the [logging config file documentation](#) for details.

Structure Examples from Diamond

Diamond is more than an executable application—it’s also a library that provides a way for users to create and use custom collectors.

We’ll highlight more things we like about the overall package structure, and then dig into how exactly Diamond makes it possible for the application to import and use externally defined collectors.

Separate different functionality into namespaces (they are one honking great idea)

The diagram in [Figure 5-2](#) shows the server module interacting with three other modules in the project: `diamond.handler`, `diamond.collector`, and `diamond.utils`. The *utils* subpackage could realistically have contained all of its classes and functions in a single, large *util.py* module, but there was an opportunity to use namespaces to separate code into related groups, and the development team took it. Honking great!

All of the implementations of Handlers are contained in *diamond/handler* (which makes sense), but the structure for the Collectors is different. There’s not a directory, only a module *diamond/collector.py* that defines the Collector and ProcessCollector.

tor base classes. All implementations of the Collectors are defined instead in *Diamond/src/collectors/* and would be installed in the virtual environment under *venv/share/diamond/collectors* when installing from PyPI (as recommended) rather than from GitHub (like we did to read it). This helps the user to create new implementations of Collectors: placing all of the collectors in the same location makes it easier for the application to find them and easier for library users to follow their example.

Finally, each Collector implementation in *Diamond/src/collectors* is in its own directory (rather than in a single file), which makes it possible to keep each Collector implementation's tests separate. Also honking great.

User-extensible custom classes (complex is better than complicated)

It's easy to add new Collector implementations: just subclass the `diamond.collector.Collector` *abstract base class*,⁶ implement a `Collector.collect()` method, and place the implementation in its own directory in *venv/src/collectors/*.

Underneath, the implementation is complex, but the user doesn't see it. This section shows both the simple user-facing part of Diamond's Collector API and the complex code that makes this user interface possible.

Complex versus complicated. We can boil down the user experience of working with *complex* code to be something like experiencing a Swiss watch—it just works, but inside there are a ton of precisely made little pieces, all interfacing with remarkable precision, in order to create the effortless user experience. Using *complicated* code, on the other hand, is like piloting an airplane—you really have to know what you're doing to not crash and burn.⁷ We don't want to live in a world without airplanes, but we *do* want our watches to work without us having to be rocket scientists. Wherever it's possible, less complicated user interfaces are a good thing.

The simple user interface. To create a custom data collector, the user must subclass the abstract class, `Collector`, and then provide, via the configuration file, the path to that new collector. Here is an example of a new Collector definition from *Diamond/src/collectors/cpu/cpu.py*. When Python searches for the `collect()` method, it will look in the `CPUCollector` for a definition first, and then if it doesn't find the definition, it

6 In Python, an abstract base class is a class that has left certain methods undefined, with the expectation that the developer will define them in the subclass. In the abstract base class, this function raises a `NotImplementedError`. A more modern alternative is to use Python's module for abstract base classes, `abc`, first implemented in Python 2.6, which will error when constructing an incomplete class rather than when trying to access that class's unimplemented method. The full specification is defined in [PEP 3119](#).

7 This is a paraphrase of a great blog post on the subject by Larry Cuban, a professor emeritus of education at Stanford, titled "[The Difference Between Complicated and Complex Matters](#)."

will use `diamond.collector.Collector.collect()`, which raises the `NotImplementedError`.

Minimal collector code would look like this:

```
# coding=utf-8
import diamond.collector
import psutil

class CPUCollector(diamond.collector.Collector):

    def collect(self):
        # In Collector, this just contains raise(NotImplementedError)
        metric_name = "cpu.percent"
        metric_value = psutil.cpu_percent()
        self.publish(metric_name, metric_value)
```

The default place to store the collector definitions is in the directory `venv/share/diamond/collectors/`; but you can store it wherever you define in the `collectors_path` value in the configuration file. The class name, `CPUCollector`, is already listed in the example configuration file. Except for adding a `hostname` or a `hostname_method` specification either in the overall defaults (under the text in the configuration file) or in the individual collector's overrides, as shown in the following example, there need not be any other changes (the documentation lists all of the optional **collector settings**):

```
[[CPUCollector]]
enabled = True
hostname_method = smart
```

The more complex internal code. Behind the scenes, the Server will call `utils.load_collectors()` using the path specified in `collectors_path`. Here is most of that function, truncated for brevity:

```
def load_collectors(paths=None, filter=None):
    """Scan for collectors to load from path"""
    # Initialize return value
    collectors = {}
    log = logging.getLogger('diamond')

    if paths is None:
        return

    if isinstance(paths, basestring): ❶
        paths = paths.split(',')
        paths = map(str.strip, paths)

    load_include_path(paths) ❷

    for path in paths:
        ##~~ Skip lines that confirm 'path' exists.
```

```

for f in os.listdir(path):

    # Are we a directory? If so, process down the tree
    fpath = os.path.join(path, f)
    if os.path.isdir(fpath):
        subcollectors = load_collectors([fpath]) ❸
        for key in subcollectors: ❹
            collectors[key] = subcollectors[key]

    # Ignore anything that isn't a .py file
    elif (os.path.isfile(fpath)
          ##~~ ... Skip tests confirming fpath is a Python module ...
          ):

        ##~~ ... Skip the part that ignores filtered paths ...
        modname = f[:-3]

    try:
        # Import the module
        mod = __import__(modname, globals(), locals(), ['*']) ❺
    except (KeyboardInterrupt, SystemExit), err:
        ##~~ ... Log the exception and quit ...
    except:
        ##~~ ... Log the exception and continue ...

    # Find all classes defined in the module
    for attrname in dir(mod):
        attr = getattr(mod, attrname) ❻
        # Only attempt to load classes that are subclasses
        # of Collectors but are not the base Collector class
        if (inspect.isclass(attr)
            and issubclass(attr, Collector)
            and attr != Collector):
            if attrname.startswith('parent_'):
                continue
            # Get class name
            fqcn = '.'.join([modname, attrname])
            try:
                # Load Collector class
                cls = load_dynamic_class(fqcn, Collector) ❼
                # Add Collector class
                collectors[cls.__name__] = cls ❽
            except Exception:
                ##~~ log the exception and continue ...

    # Return Collector classes
    return collectors

```

- ❶ Break up the string (first function call); otherwise, the paths are lists of string paths to where the user-defined custom Collector subclasses are defined.

- ② This recursively descends the paths given, inserting every directory into `sys.path` so that later the Collectors can be imported.
- ③ Here's the recursion—`load_collectors()` is calling itself.⁸
- ④ After loading the subdirectories' collectors, update the original dictionary of custom collectors with the new ones from those subdirectories.
- ⑤ Since the introduction of Python 3.1, the `importlib` module in Python's standard library provides a preferred way to do this (via the module `importlib.import_module`; parts of `importlib.import_module` have also been backported to Python 2.7). This demonstrates how to programmatically import a module given the string module name.
- ⑥ Here's how to programmatically access attributes in a module given just the string attribute name.
- ⑦ Actually, `load_dynamic_class` may not be necessary here. It re-imports the module, checks that the named class is actually a class, checks that it's actually a Collector, and if so returns the newly loaded class. Redundancies sometimes occur in open source code written by large groups.
- ⑧ Here's how they get the class name to use later when applying the configuration file options given only the string class name.

Style Examples from Diamond

There's a great example use of a closure in Diamond that demonstrates what was said in “[Late binding closures](#)” on page 59 about this behavior often being desirable.

Example use of a closure (when the gotcha isn't a gotcha)

A *closure* is a function that makes use of variables available in the local scope that would otherwise not be available when the function is called. They can be difficult to implement and understand in other languages, but are not hard to implement in Python, because Python treats functions just like any other object.⁹ For example, functions can be passed around as arguments, or returned from other functions.

⁸ Python has a recursion limit (a maximum number of times a function is allowed to call itself) that's relatively restrictive by default, to discourage excessive use of recursion. Get your recursion limit by typing `import sys; sys.getrecursionlimit()`.

⁹ Programming languages that can do this are said to have *first-class functions*—functions are treated as first-class citizens, like any other object.

Here's an example excerpt from the `diamond` executable that shows how to implement a closure in Python:

```
##~ ... Skip the import statements ... ❶

def main():
    try:
        ##~ ... Skip code that creates the command-line parser ...

        # Parse command-line Args
        (options, args) = parser.parse_args()

        ##~ ... Skip code that parses the configuration file ...
        ##~ ... Skip code that sets up the logger ...

        # Pass the exit upstream rather than handle it as an general exception
        except SystemExit, e:
            raise SystemExit

        ##~ ... Skip code that handles other exceptions related to setup ...

    try:
        # PID MANAGEMENT ❷
        if not options.skip_pidfile:
            # Initialize PID file
            if not options.pidfile:
                options.pidfile = str(config['server']['pid_file'])

            ##~ ... Skip code to open and read the PID file if it exists, ...
            ##~ ... and then delete the file if there is no such PID ...
            ##~ ... or exits if there is already a running process. ...

            ##~ ... Skip the code that sets the group and user ID ...
            ##~ ... and the code that changes the PID file permissions. ...

            ##~ ... Skip the code that checks whether to run as a daemon, ...
            ##~ ... and if so detaches the process. ...

        # PID MANAGEMENT ❸
        if not options.skip_pidfile:
            # Finish initializing PID file
            if not options.foreground and not options.collector:
                # Write PID file
                pid = str(os.getpid())
                try:
                    pf = file(options.pidfile, 'w+')
                except IOError, e:
                    log.error("Failed to write child PID file: %s" % (e))
                    sys.exit(1)
                pf.write("%s\n" % pid)
```

```

        pf.close()
        # Log
        log.debug("Wrote child PID file: %s" % (options.pidfile))

    # Initialize server
    server = Server(configfile=options.configfile)

    def sigint_handler(signum, frame): ❷
        log.info("Signal Received: %d" % (signum))
        # Delete PID file
        if not options.skip_pidfile and os.path.exists(options.pidfile): ❸
            os.remove(options.pidfile)
            # Log
            log.debug("Removed PID file: %s" % (options.pidfile))
            sys.exit(0)

    # Set the signal handlers
    signal.signal(signal.SIGINT, sigint_handler) ❹
    signal.signal(signal.SIGTERM, sigint_handler)

    server.run()

    # Pass the exit upstream rather than handle it as a general exception
    except SystemExit, e:
        raise SystemExit

    ##~~ ... Skip code that handles any other exceptions ...
    ##~~ ... and all of the rest of the script.

```

- ❶ When we skip code, the missing parts will be summarized by a comment preceded by two tildes (`##~~` like this).
- ❷ The reason for the PID¹⁰ file is to make sure the daemon is unique (i.e., not accidentally started twice), to communicate the associated process ID quickly to other scripts, and to make it evident that an abnormal termination has occurred (because in this script, the PID file is deleted upon normal termination).
- ❸ All of this code is just to provide context leading up to the closure. At this point, either the process is running as a daemon (and now has a different process ID than before) or it will skip this part because it's already written its correct PID to the PID file.

¹⁰ PID stands for “process identifier.” Every process has a unique identifier that is available in Python using the `os` module in the standard library: `os.getpid()`.

- ④ This (`sigint_handler()`) is the closure. It is defined inside of `main()`, rather than at the top level, outside of any functions, because it needs to know whether to look for a PID file, and if so where to look.
- ⑤ It gets this information from the command-line options, which it can't obtain until after the call to `main()`. That means all of the options related to the PID file are local variables in `main`'s namespace.
- ⑥ The closure (the function `sigint_handler()`) is sent to the signal handler and will be used to handle `SIGINT` and `SIGTERM`.

Tablib

Tablib is a Python library that converts between different data formats, storing data in a `Dataset` object, or multiple `Datasets` in a `Databook`. `Datasets` stored in the JSON, YAML, DBF, and CSV file formats can be imported, and `datasets` can be exported to XLSX, XLS, ODS, JSON, YAML, DBF, CSV, TSV, and HTML. Tablib was first released by Kenneth Reitz in 2010. It has the intuitive API design typical of Reitz's projects.

Reading a Small Library

Tablib is a library, not an application, so there isn't a single obvious entry point like there is with `HowDoI` and `Diamond`.

Get Tablib from GitHub:

```
$ git clone https://github.com/kennethreitz/tablib.git
$ virtualenv -p python3 venv
$ source venv/bin/activate
(venv)$ cd tablib
(venv)$ pip install --editable .
(venv)$ python test_tablib.py # Run the unit tests.
```

Read Tablib's documentation

Tablib's documentation starts off immediately with a use case, and then goes into describing its capabilities in more detail: it provides a `Dataset` object that has rows, headers, and columns. You can do I/O from various formats to the `Dataset` object. And the advanced usage section says you can add tags to rows, and create derived columns that are functions of other columns.

Use Tablib

Tablib is a library, not an executable like `HowDoI` or `Diamond`, so you can open a Python interactive session and have the expectation that you can use the `help()`

function to explore the API. Here's our example of the `tablib.Dataset` class, the different data formats, and how I/O works:

```
>>> import tablib
>>> data = tablib.Dataset()
>>> names = ('Black Knight', 'Killer Rabbit')
>>>
>>> for name in names:
...     fname, lname = name.split()
...     data.append((fname, lname))
...
>>> data.dict
[['Black', 'Knight'], ['Killer', 'Rabbit']]
>>>
>>> print(data.csv)
Black,Knight
Killer,Rabbit

>>> data.headers=('First name', 'Last name')
>>> print(data.yaml)
- {First name: Black, Last name: Knight}
- {First name: Killer, Last name: Rabbit}

>>> with open('tmp.csv', 'w') as outfile:
...     outfile.write(data.csv)
...
64
>>> newdata = tablib.Dataset()
>>> newdata.csv = open('tmp.csv').read()
>>> print(newdata.yaml)
- {First name: Black, Last name: Knight}
- {First name: Killer, Last name: Rabbit}
```

Read Tablib's code

The file structure under *tablib/* looks like this:

```
tablib
|--- __init__.py
|--- compat.py
|--- core.py
|--- formats/
|--- packages/
```

The two directories, *tablib/formats/* and *tablib/packages/*, will be discussed in a few sections.

Python supports module-level docstrings as well as the docstrings we've already described—a string literal that is the first statement in a function, class, or class method. Stack Overflow has good advice on how to [document a module](#). For us, this means another way to explore source code is by typing `head *.py` in a terminal shell

while in the directory at the top level of the package—to show all of the module docstrings at once. Here’s what we get:

```
(venv)$ cd tablib
(venv)$ head *.py
==> __init__.py <== ❶
""" Tablib. """

from tablib.core import (
    Databook, Dataset, detect, import_set, import_book,
    InvalidDatasetType, InvalidDimensions, UnsupportedFormat,
    __version__
)

==> compat.py <== ❷
# -*- coding: utf-8 -*-

"""
tablib.compat
~~~~~

Tablib compatibility module.

"""

==> core.py <== ❸
# -*- coding: utf-8 -*-
"""
    tablib.core
    ~~~~~

    This module implements the central Tablib objects.

    :copyright: (c) 2014 by Kenneth Reitz.
    :license: MIT, see LICENSE for more details.
"""
```

We learn that:

- ❶ The top-level API (the contents of `__init__.py` are accessible from `tablib` after an `import tablib` statement) has just nine entry points: the `Databook` and `Dataset` classes are mentioned in the documentation, `detect` could be for identifying formatting, `import_set` and `import_book` must import data, and the last three classes—`InvalidDatasetType`, `InvalidDimensions`, and `UnsupportedFormat`—look like exceptions. (When code follows PEP 8, we can tell which objects are custom classes from their capitalization.)

- ② *tablib/compat.py* is a compatibility module. A quick look inside will show that it handles Python 2/Python 3 compatibility issues in a similar way to HowDoI, by resolving different locations and names to the same symbol for use in *tablib/core.py*.
- ③ *tablib/core.py*, like it says, implements the central Tablib objects like Dataset and Databook.

Tablib's Sphinx Documentation

Tablib's **documentation** provides a good example use of **Sphinx** because it's a small library, and it makes use of a lot of **Sphinx extensions**.

The documentation's current Sphinx build is at **Tablib's documentation page**. If you want to build the documentation yourself (Windows users will need a **make command**—it's old but works fine), do this:

```
(venv)$ pip install sphinx
(venv)$ cd docs
(venv)$ make html
(venv)$ open _build/html/index.html # To view the result.
```

Sphinx provides a number of **theme options** with default layout templates and CSS themes. Tablib's templates for two of the notes on the left sidebar are in *docs/_templates/*. Their names are not arbitrary; they're in *basic/layout.html*. You can find that file in the Sphinx themes directory, which can be located by typing this on the command line:

```
(venv)$ python -c 'import sphinx.themes;print(sphinx.themes.__path__)'
```

Advanced users can also look in *docs/_themes/kr/*, a custom theme that extends the basic layout. It is selected by adding the *_themes/* directory to the system path, setting `html_theme_path = ['_themes']` and setting `html_theme = 'kr'` in *docs/conf.py*.

To include API documentation that's automatically generated from the docstrings in your code, use `autoclass::`. You have to copy the docstring formatting in Tablib for this to work:

```
.. autoclass:: Dataset
   :inherited-members:
```

To get this functionality, you have to answer “yes” to the question about including the “autodoc” Sphinx extension when you run `sphinx-quickstart` to create a new Sphinx project. The `:inherited-members:` directive also adds documentation for the attributes inherited from parent classes.

Structure Examples from Tablib

The primary thing we want to highlight from Tablib is the absence of the use of classes in the modules in *tablib/formats/*—it’s a perfect example of the statement we made earlier about not overusing classes. Next, we show excerpts of how Tablib uses the decorator syntax and the **property class** to create derived attributes like the dataset’s height and width, and how it dynamically registers file formats to avoid duplicating what would be boilerplate code for each of the different format types (CSV, YAML, etc.).

The last two subsections are a little obscure—we look at how Tablib vendorizes dependencies, and then discuss the `__slots__` property of new class objects. You can skip these sections and still lead a happy, Pythonic life.

No needless object-oriented code in formats (use namespaces for grouping functions)

The *formats* directory contains all of the defined file formats for I/O. The module names, *_csv.py*, *_tsv.py*, *_json.py*, *_yaml.py*, *_xls.py*, *_xlsx.py*, *_ods.py*, and *_xls.py* are prefixed with an underscore—this indicates to the library user that they are not intended for direct use. We can change directories into *formats*, and search for classes and functions. Using `grep ^class formats/*.py` reveals there are no class definitions, and using `grep ^def formats/*.py` shows that each module contains some or all of the following functions:

- `detect(stream)` infers the file format based on the stream content.
- `dset_sheet(dataset, ws)` formats the Excel spreadsheet cells.
- `export_set(dataset)` exports the Dataset to the given format, returning a formatted string with the new format. (Or, for Excel, returning a bytes object—or a binary-formatted string in Python 2.)
- `import_set(dset, in_stream, headers=True)` replaces the contents of the dataset with the contents of the input stream.
- `export_book(databook)` exports the Datasheets in the Databook to the given format, returning a string or bytes object.
- `import_book(dbook, in_stream, headers=True)` replaces the contents of the databook with the contents of the input stream.

This is an example of using modules as namespaces (after all, they *are* one honking great idea) to separate functions, rather than using unnecessary classes. We know each function’s purpose from its name: for example, `formats._csv.import_set()`, `formats._tsv.import_set()`, and `formats._json.import_set()` import datasets from CSV, TSV, and JSON-formatted files, respectively. The other functions do data

exporting and file format detection, when possible, for each of Tablib’s available formats.

Descriptors and the property decorator (engineer immutability when the API would benefit)

Tablib is our first library that uses Python’s decorator syntax, described in “Decorators” on page 67. The syntax uses the `@` symbol in front of a function name, placed directly above another function. It modifies (or “decorates”) the function directly below. In the following excerpt, property changes the functions `Dataset.height` and `Dataset.width` into descriptors—classes with at least one of the `__get__()`, `__set__()`, or `__delete__()` (“getter”, “setter”, or “delete”) methods defined. For example, the attribute lookup `Dataset.height` will trigger the getter, setter, or delete function depending on the context in which that attribute is used. This behavior is only possible for new-style classes, discussed momentarily. See this useful Python [tutorial on descriptors](#) for more information.

```
class Dataset(object):
    #
    # ... omit the rest of the class definition for clarity
    #

    @property ❶
    def height(self):
        """The number of rows currently in the :class:`Dataset`.
           Cannot be directly modified. ❷
        """
        return len(self._data)

    @property
    def width(self):
        """The number of columns currently in the :class:`Dataset`.
           Cannot be directly modified.
        """
        try:
            return len(self._data[0])
        except IndexError:
            try:
                return len(self.headers)
            except TypeError:
                return 0
```

- ❶ This is how to use a decorator. In this case, property modifies `Dataset.height` to behave as a property rather than as a bound method. It can only operate on class methods.

- ② When property is applied as a decorator, the height attribute will return the height of the Dataset but it is not possible to assign a height to the Dataset by invoking `Dataset.height`.

Here is what the height and width attributes look like when used:

```
>>> import tablib
>>> data = tablib.Dataset()
>>> data.header = ("amount", "ingredient")
>>> data.append(("2 cubes", "Arcturan Mega-gin"))
>>> data.width
2
>>> data.height
1
>>>
>>> data.height = 3
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: can't set attribute
```

So, `data.height` can be accessed like an attribute, but it's not settable—it's calculated from the data and so is always current. This is ergonomic API design: `data.height` is easier to type than `data.get_height()`; it's clear what the meaning of `data.height` is; and because it is calculated from the data (and the property is not settable—only the “getter” function is defined), there isn't a danger that it will be out of sync from the correct number.

The property decorator can only be applied to attributes of classes, and only to classes that derive from the base object object (e.g., `class MyClass(object)` not `class MyClass()`—inheritance from `object` is always the case in Python 3).

This same tool is used to create Tablib's data import and export API in the various formats: Tablib does not store the string value for each of the CSV, JSON, and YAML outputs. Rather, the Dataset attributes `csv`, `json`, and `yaml` are properties, like `Data.set.height` and `Dataset.width` in the preceding example—they call a function that generates the result from the stored data or parses the input format and then replaces the core data. But there's only one dataset.

When `data.csv` is on the left of an equals sign, the property's “setter” function is called, to parse the dataset from the CSV format. And when `data.yaml` is on the right of an equals sign, or alone, the “getter” is called, to create a string with the given format from the internal dataset. Here is an example:

```
>>> import tablib
>>> data = tablib.Dataset()
>>>
>>> data.csv = "\n".join(( ①
...     "amount,ingredient",
...     "1 bottle,0l' Janx Spirit",
```

```

...     "1 measure,Santraginus V seawater",
...     "2 cubes,Arcturan Mega-gin",
...     "4 litres,Fallian marsh gas",
...     "1 measure,Qalactin Hypermint extract",
...     "1 tooth,Algolian Suntiger",
...     "a sprinkle,Zamphuur",
...     "1 whole,olive"))
>>>
>>> data[2:4]
[('2 cubes', 'Arcturan Mega-gin'), ('4 litres', 'Fallian marsh gas')]
>>>
>>> print(data.yaml) ❷
- {amount: 1 bottle, ingredient: Ol Janx Spirit}
- {amount: 1 measure, ingredient: Santraginus V seawater}
- {amount: 2 cubes, ingredient: Arcturan Mega-gin}
- {amount: 4 litres, ingredient: Fallian marsh gas}
- {amount: 1 measure, ingredient: Qalactin Hypermint extract}
- {amount: 1 tooth, ingredient: Algolian Suntiger}
- {amount: a sprinkle, ingredient: Zamphuur}
- {amount: 1 whole, ingredient: olive}

```

- ❶ `data.csv` on the lefthand side of the equals sign (assignment operator) invokes `formats.csv.import_set()`, with `data` as the first argument, and the string of Gargle Blaster ingredients as its second argument.
- ❷ `data.yaml` alone invokes `formats.yaml.export_set()`, with `data` as its argument, outputting the formatted YAML string for the `print()` function.

The “getter”, “setter”, and also a “deleter” function can be bound to a single attribute using `property`. Its signature is `property(fget=None, fset=None, fdel=None, doc=None)`, in which `fget` identifies the “getter” function (`formats.csv.import_set()`), `fset` identifies the “setter” function (`formats.csv.export_set()`), and `fdel` identifies the “deleter” function, which is left as `None`. We will see the code where the formatting properties are set, programmatically, next.

Programmatically registered file formats (don’t repeat yourself)

Tablib places all of the file formatting routines in the *formats* subpackage. This structure choice makes the main *core.py* module cleaner and the entire package modular; it’s easy to add new file formats. Although it would have been possible to paste chunks of nearly identical code and import each file format’s import and export behaviors separately, all of the formats are *programmatically* loaded into the *Dataset* class to properties named after each format.

We're printing the entire contents of *formats/__init__.py* in the following code example because it's not too large a file, and we want to show where *formats.available* is defined:

```
# -*- coding: utf-8 -*- ❶

""" Tablib - formats
"""

from . import _csv as csv
from . import _json as json
from . import _xls as xls
from . import _yaml as yaml
from . import _tsv as tsv
from . import _html as html
from . import _xlsx as xlsx
from . import _ods as ods

available = (json, xls, yaml, csv, tsv, html, xlsx, ods) ❷
```

- ❶ This line explicitly tells the Python interpreter that the file encoding is UTF-8.¹¹
- ❷ Here's the definition of *formats.available*, right in *formats/__init__.py*. It's also available via *dir(tablib.formats)*, but this explicit list is easier to understand.

In *core.py*, rather than about 20 (ugly, hard to maintain) repeated function definitions for each format option, the code imports each format programmatically by calling *self._register_formats()* at the end of the *Dataset*'s *__init__()* method. We've excerpted just *Dataset._register_formats()* here:

```
class Dataset(object):
    #
    # ... skip documentation and some definitions ...
    #

    @classmethod ❶
    def _register_formats(cls):
        """Adds format properties."""
        for fmt in formats.available: ❷
            try:
                setattr(cls, fmt.title,
                        property(fmt.export_set, fmt.import_set)) ❸
            except AttributeError: ❹
                setattr(cls, fmt.title, property(fmt.export_set)) ❺
```

11 ASCII is default in Python 2, and UTF-8 is default in Python 3. There are multiple allowed ways to communicate encoding, all listed in [PEP 263](#). You can use the one that works best with your favorite text editor.

```

except AttributeError:
    pass ❹

#
# ... skip more definitions ...
#

@property ❺
def tsv():
    """A TSV representation of the :class:`Dataset` object. The top
    row will contain headers, if they have been set. Otherwise, the
    top row will contain the first row of the dataset.

    A dataset object can also be imported by setting
    the :class:`Dataset.tsv` attribute. ::

        data = tablib.Dataset()
        data.tsv = 'age|tfirst_name|tlast_name|n90|tJohn|tAdams' ❸

    Import assumes (for now) that headers exist.
    """
    pass

```

- ❶ The @classmethod symbol is a *decorator*, described more extensively in “Decorators” on page 67, that modifies the method `_register_formats()` so that it passes the object’s class (`Dataset`) rather than the object instance (`self`) as its first argument.
- ❷ The `formats.available` is defined in `formats/__init__.py` and contains all of the available formatting options.
- ❸ In this line, `setattr` assigns a value to the attribute named `fmt.title` (i.e., `Dataset.csv` or `Dataset.xls`). The value it assigns is a special one; `property(fmt.export_set, fmt.import_set)` turns `Dataset.csv` into a *property*.
- ❹ There will be an `AttributeError` if `fmt.import_set` is not defined.
- ❺ If there is no import function, try to assign just the export behavior.
- ❻ If there is neither an export nor an import function to assign, just don’t assign anything.
- ❼ Each of the file formats is defined as a property here, with a descriptive docstring. The docstring will be retained when `property()` is called at tag ❸ or ❺ to assign the extra behaviors.

- 8 The `\t` and `\n` are string escape sequences that represent the Tab character and a newline, respectively. They're all listed in Python's [string literals documentation](#).

But We Are All Responsible Users

These uses of the `@property` decorator are *not* like the uses of similar tools in Java, where the goal is to control the user's access to data. That goes against the Python philosophy that *we are all responsible users*. Instead, the purpose of `@property` is to separate the data from view functions related to the data (in this case, the height, width, and various storage formats). When there doesn't need to be a preprocessing or post-processing “getter” or “setter” function, the more Pythonic option is to just assign the data to a regular attribute and let the user interact with it.

Vendorized dependencies in packages (an example of how to vendorize)

Tablib's dependencies are currently vendorized (meaning they are shipped bundled with the code—in this case, in the directory *packages*) but may be moved to a plug-in system in the future. The *packages* directory contains third-party packages included inside Tablib to ensure compatibility, rather than the other option, which is to specify versions in the *setup.py* file that will be downloaded and installed when Tablib is installed. This technique is discussed in “[Vendorizing Dependencies](#)” on page 71; the choice for Tablib was made both to reduce the number of dependencies the user would have to download, and because sometimes there are different packages for Python 2 and Python 3, which are both included. (The appropriate one is imported, and their functions set to a common name, in *tablib/compat.py*). That way, Tablib can have one single code base instead of two—one for each version of Python. Because each of these dependencies has its own license, a *NOTICE* document was added to the top level of the project directory that lists each dependency's license.

Saving memory with `__slots__` (optimize judiciously)

Python prefers readability over speed. Its entire design, its Zen aphorisms, and its early [influence from educational languages like ABC](#) are all about placing user-friendliness above performance (although we'll talk about more optimization options in “[Speed](#)” on page 223).

The use of `__slots__` in tablib is a case where optimization matters. This is a slightly obscure reference, and it's only available for new-style classes (described in a few pages), but we want to show that it's possible to optimize Python when necessary. This optimization is only useful when you have tons of very small objects by reducing the footprint of each class instance by the size of one dictionary (large objects would make this small savings irrelevant, and fewer objects make the savings not worth it). Here is an excerpt from the [__slots__ documentation](#):

By default, instances of classes have a dictionary for attribute storage. This wastes space for objects having very few instance variables. The space consumption can become acute when creating large numbers of instances.

The default can be overridden by defining `__slots__` in a class definition. The `__slots__` declaration takes a sequence of instance variables and reserves just enough space in each instance to hold a value for each variable. Space is saved because `__dict__` is not created for each instance.

Normally, this isn't something to care about—notice that `__slots__` doesn't appear in the `Dataset` or `Databook` classes, just the `Row` class—but because there can be thousands of rows of data, `__slots__` is a good idea. The `Row` class is not exposed in `tablib/__init__.py` because it is a helper class to `Dataset`, instantiated once for every row. This is how its definition looks in the beginning part of the definition of the `Row` class:

```
class Row(object):
    """Internal Row object. Mainly used for filtering."""

    __slots__ = ['_row', 'tags']

    def __init__(self, row=list(), tags=list()):
        self._row = list(row)
        self.tags = list(tags)

    #
    # ... etc. ...
    #
```

The problem now is that there is no longer a `__dict__` attribute in the `Row` instances, but the `pickle.dump()` function (used for object serialization) by default uses `__dict__` to serialize the object unless the method `__getstate__()` is defined. Likewise, during unpickling (the process that reads the serialized bytes and reconstructs the object in memory), if `__setstate__()` is not defined, `pickle.load()` will load to the object's `__dict__` attribute. Here is how to get around that:

```
class Row(object):
    #
    # ... skip the other definitions ...
    #

    def __getstate__(self):

        slots = dict()

        for slot in self.__slots__:
            attribute = getattr(self, slot)
            slots[slot] = attribute
        return slots

    def __setstate__(self, state):
```

```

for (k, v) in list(state.items()):
    setattr(self, k, v)

```

For more information about `__getstate__()` and `__setstate__()` and pickling, see the [__getstate__ documentation](#).

Style Examples from Tablib

We have one single style example from Tablib—operator overloading—which gets into the details of Python’s data model. Customizing the behavior of your classes makes it easier for those who use your API to write beautiful code.

Operator overloading (beautiful is better than ugly)

This code section uses Python’s operator overloading to enable operations on either the Dataset’s rows or columns. The following first sample code shows interactive use of the bracket operator (`[]`) for both numerical indices and column names, and the second one shows the code that uses this behavior:

```

>>> data[-1] ❶
('1 whole', 'olive')
>>>
>>> data[-1] = ['2 whole', 'olives'] ❷
>>>
>>> data[-1]
('2 whole', 'olives') ❸
>>>
>>> del data[2:7] ❹
>>>
>>> print(data.csv)
amount,ingredient ❺
1 bottle,0l' Janx Spirit
1 measure,Santraginus V seawater
2 whole,olives

>>> data['ingredient'] ❻
["0l' Janx Spirit", 'Santraginus V seawater', 'olives']

```

- ❶ When using numbers, accessing data via the bracket operator (`[]`) gives the row at the specified location.
- ❷ This is assignment using the bracket operator ...
- ❸ ... it becomes 2 olives instead of the original one.
- ❹ This is deletion using a *slice*—`2:7` denotes all of the numbers 2,3,4,5,6 but not 7.

- ⑤ See how the recipe afterward is much smaller.
- ⑥ It is also possible to access columns by name.

The part of the `Dataset` code that defines the behavior of the bracket operator shows how to handle access both by column name and by row number:

```
class Dataset(object):
    #
    # ... skip the rest of the definitions for brevity ...
    #

    def __getitem__(self, key):
        if isinstance(key, str) or isinstance(key, unicode): ①
            if key in self.headers: ②
                pos = self.headers.index(key) # get 'key' index from each data
                return [row[pos] for row in self._data]
            else: ③
                raise KeyError
        else:
            _results = self._data[key]
            if isinstance(_results, Row): ④
                return _results.tuple
            else:
                return [result.tuple for result in _results] ⑤

    def __setitem__(self, key, value): ⑥
        self._validate(value)
        self._data[key] = Row(value)

    def __delitem__(self, key):
        if isinstance(key, str) or isinstance(key, unicode): ⑦
            if key in self.headers:
                pos = self.headers.index(key)
                del self.headers[pos]

                for row in self._data:
                    del row[pos]
            else:
                raise KeyError
        else:
            del self._data[key]
```

- ① First, check whether we are seeking a column (True if key is a string) or a row (True if the key is an integer or slice).
- ② Here the code checks for the key to be in `self.headers`, and then...

- ③ ...explicitly raises a `KeyError` so that access by column name behaves as one would expect a dictionary to. The whole `if/else` pair is not necessary for the operation of the function—if it were omitted, a `ValueError` would still be raised by `self.headers.index(key)` if `key` were not in `self.headers`. The only purpose for this check is to provide a more informative error for the library user.
- ④ This is how the code determines whether `key` was a number or a slice (like `2:7`). If a slice, the `_results` would be a list, not a `Row`.
- ⑤ Here is where the slice is processed. Because the rows are returned as tuples, the values are an immutable copy of the actual data, and the dataset's values (actually stored as lists) won't accidentally be corrupted by an assignment.
- ⑥ The `__setitem__()` method can change a single row but not a column. This is intentional; there is no way provided to change the content of an entire column; and for data integrity, this is probably not a bad choice. The user can always transform the column and insert it at any position using one of the methods `insert_col()`, `lpush_col()`, or `rpush_col()`.
- ⑦ The `__delitem__()` method can either delete a column or a row, using the same logic as `__getitem__()`.

For more information about additional operator overloading and other special methods, see the Python documentation on [Special method names](#).

Requests

On Valentine's day in 2011, Kenneth Reitz released a love letter to the Python community: the Requests library. Its enthusiastic adoption emphatically makes the case for intuitive API design (meaning the API is so straightforward you almost don't need documentation).

Reading a Larger Library

Requests is a larger library than Tablib, with many more modules, but we'll still approach reading it the same way—by looking at the documentation and following the API through the code.

Get Requests from GitHub:

```
$ git clone https://github.com/kennethreitz/requests.git
$ virtualenv -p python3 venv
$ source venv/bin/activate
(venv)$ cd requests
(venv)$ pip install --editable .
```

```
(venv)$ pip install -r requirements.txt # Required for unit tests
(venv)$ py.test tests # Run the unit tests.
```

Some tests may fail—for example, if your service provider intercepts 404 errors to give some advertising page, you won't get the `ConnectionError`.

Read Requests's documentation

Requests is a bigger package, so first just scan the section titles from the [Requests documentation](#). Requests extends `urllib` and `httplib` from Python's standard library to provide methods that perform HTTP requests. The library includes support for international domains and URLs, automatic decompression, automatic content decoding, browser style SSL verification, HTTP(S) proxy support, and other features, which are all defined by the Internet Engineering Task Force (IETF) standards for HTTP in their requests for comment (RFCs) 7230 through 7235.¹²

Requests strives to cover *all* of the IETF's HTTP specifications, using only a handful of functions, a bunch of keyword arguments, and a few featureful classes.

Use Requests

Like with Tablib, there is enough information in the docstrings to use Requests without actually reading the online documentation. Here's a brief interaction:

```
>>> import requests
>>> help(requests) # Shows a usage statement and says to see `requests.api`
>>> help(requests.api) # Shows a detailed API description
>>>
>>> result = requests.get('https://pypi.python.org/pypi/requests/json')
>>> result.status_code
200
>>> result.ok
True
>>> result.text[:42]
'{"\n    "info": {\n        "maintainer": null'
>>>
>>> result.json().keys()
dict_keys(['info', 'releases', 'urls'])
>>>
>>> result.json()['info']['summary']
'Python HTTP for Humans.'
```

Read Requests's code

Here are the contents of the Requests package:

¹² If you need a vocabulary refresh, RFC 7231 is the [HTTP semantics document](#). If you scan the table of contents and read the introduction, you'll know enough about the scope to tell whether the definition you want is covered, and where to find it.

```
$ ls
__init__.py      cacert.pem ❶ exceptions.py  sessions.py
adapters.py      certs.py    hooks.py      status_codes.py
api.py           compat.py   models.py     structures.py
auth.py          cookies.py  packages/ ❷  utils.py
```

- ❶ *cacert.pem* is a default certificate bundle to use when checking SSL certificates.
- ❷ Requests has a flat structure, except for a *packages* directory that vendorizes (contains the external libraries) *chardet* and *urllib3*. These dependencies are imported as `requests.packages.chardet` and `requests.packages.urllib3`, so programmers can still access *chardet* and *urllib3* from the standard library.

We can mostly figure out what’s happening thanks to well-chosen module names, but if we want a little more information, we can again peek at the module docstrings by typing `head *.py` in the top-level directory. The following lists displays these module docstrings, slightly truncated. (It doesn’t show *compat.py*. We can tell from its name, especially because it’s named the same as in Reitz’s *Tablib* library, that it takes care of Python 2 to Python 3 compatibility.)

api.py

Implements the Requests API

hooks.py

Provides the capabilities for the Requests hooks system

models.py

Contains the primary objects that power Requests

sessions.py

Provides a Session object to manage and persist settings across requests (cookies, auth, proxies)

auth.py

Contains the authentication handlers for Requests

status_codes.py

A lookup table mapping status titles to status codes

cookies.py

Compatibility code to be able to use `cookielib.CookieJar` with requests

adapters.py

Contains the transport adapters Requests uses to define and maintain connections

exceptions.py

All of Requests’ exceptions

structures.py

Data structures that power Requests

certs.py

Returns the preferred default CA certificate bundle listing trusted SSL certificates

utils.py

Provides utility functions that are used within Requests that are also useful for external consumption

Insights from reading all of the headers:

- There is a hook system (*hooks.py*), implying the user can modify how Requests works. We won't discuss it in depth because it will take us too far off topic.
- The main module is *models.py*, as it contains “the primary objects that power Requests.”
- The reason `sessions.Session` exists is to persist cookies across multiple requests (that might occur during authentication, for example).
- The actual HTTP connection is made by objects from *adapters.py*.
- The rest is kind of obvious: *auth.py* is for authentication, *status_codes.py* has the status codes, *cookies.py* is for adding and removing cookies, *exceptions.py* is for exceptions, *structures.py* contains data structures (e.g., a case-insensitive dictionary), and *utils.py* contains utility functions.

The idea to put communication separately in *adapters.py* is innovative (at least to this writer). It means `models.Request`, `models.PreparedRequest`, and `models.Response` don't actually *do* anything—they *just store data*, possibly manipulating it a bit for presentation, pickling, or encoding purposes. Actions are handled by separate classes that exist specifically to perform an action, like authentication or communication. Every class does *just one thing*, and each module contains classes that do similar things—a Pythonic approach most of us already adhere to with our function definitions.

Requests's Sphinx-Compatible Docstrings

If you are starting a new project and using Sphinx and its autodoc extension, you will need to format your docstrings so that Sphinx can parse them.

The Sphinx documentation is not always easy to search for what keywords to place where. Many people actually recommend copying the docstrings in Requests if you want to get the format right, rather than find the instructions in the Sphinx docs. For example, here is the definition of `delete()` in `requests/api.py`:

```
def delete(url, **kwargs):
    """Sends a DELETE request.

    :param url: URL for the new :class:`Request` object.
    :param |*|*kwargs: Optional arguments that ``request`` takes.
    :return: :class:`Response <Response>` object
    :rtype: requests.Response
    """

    return request('delete', url, **kwargs)
```

The Sphinx autodoc rendering of this definition is in the [online API documentation](#).

Structure Examples from Requests

Everyone loves the Requests API—it is easy to remember and helps its users to write simple, beautiful code. This section first discusses the design preference for more comprehensible error messages and an easy-to-memorize API that we think went into creation of the `requests.api` module, and then explores the differences between the `requests.Request` and `urllib.request.Request` object, offering an opinion on why `requests.Request` exists.

Top-level API (preferably only one obvious way to do it)

The functions defined in `api.py` (except `request()`) are named after HTTP request methods.¹³ Each request method is the same except for its method name and the choice of exposed keyword parameters, so we're truncating this excerpt from `requests/api.py` after the `get()` function:

```
# -*- coding: utf-8 -*-

"""
requests.api
~~~~~
```

¹³ These are defined in [section 4.3 of the current Hypertext Transfer Protocol request for comments](#).

This module implements the Requests API.

:copyright: (c) 2012 by Kenneth Reitz.

:license: Apache2, see LICENSE for more details.

"""

```
from . import sessions
```

```
def request(method, url, **kwargs): ❶
    """Constructs and sends a :class:`Request <Request>`.

    :param method: method for the new :class:`Request` object.
    :param url: URL for the new :class:`Request` object.
    :param params: (optional) Dictionary or bytes to be sent in the query string
                    for the :class:`Request`.

    ... skip the documentation for the remaining keyword arguments ... ❷

    :return: :class:`Response <Response>` object
    :rtype: requests.Response

    Usage::

    >>> import requests
    >>> req = requests.request('GET', 'http://httpbin.org/get')
    <Response [200]>
    """

    # By using the 'with' statement, we are sure the session is closed, thus we
    # avoid leaving sockets open which can trigger a ResourceWarning in some
    # cases, and look like a memory leak in others.
    with sessions.Session() as session: ❸
        return session.request(method=method, url=url, **kwargs)

def get(url, params=None, **kwargs): ❹
    """Sends a GET request.

    :param url: URL for the new :class:`Request` object.
    :param params: (optional) Dictionary or bytes to be sent in the query string
                    for the :class:`Request`.
    :param |*|kwargs: Optional arguments that ``request`` takes.
    :return: :class:`Response <Response>` object
    :rtype: requests.Response
    """

    kwargs.setdefault('allow_redirects', True) ❺
    return request('get', url, params=params, **kwargs) ❻
```

- ❶ The `request()` function contains a `**kwargs` in its signature. This means extra-neous keyword arguments will not cause an exception, and it hides options from the user.
- ❷ The documentation omitted here for brevity describes every keyword argument that has an associated action. If you use `**kwargs` in your function signature, this is the only way the user can tell what the contents of `**kwargs` should be, short of looking at the code themselves.
- ❸ The `with` statement is how Python supports a runtime context. It can be used with any object that has an `__enter__()` and an `__exit__()` method defined. `__enter__()` will be called upon entering the `with` statement, and `__exit__()` will be called upon exit, regardless of whether that exit is normal or due to an exception.
- ❹ The `get()` function specifically pulls out the `params=None` keyword, applying a default value of `None`. The `params` keyword argument relevant for `get` because it's for terms to be used in an HTTP query string. Exposing selected keyword arguments gives flexibility to the advanced user (via the remaining `**kwargs`) while making usage obvious for the 99% of people who don't need advanced options.
- ❺ The default for the `request()` function is to not allow redirects, so this step sets it to `True` unless the user set it already.
- ❻ The `get()` function then simply calls `request()` with its first parameter set to `"get"`. Making `get` a function has two advantages over just using a string argument like `request("get", ...)`. First, it becomes obvious, even without documentation, which HTTP methods are available with this API. Second, if the user makes a typographical error in the method name, a `NameError` will be raised earlier, and probably with a less confusing traceback, than would happen with error checking deeper in the code.

No new functionality is added in *requests/api.py*; it exists to present a simple API for the user. Plus, putting the string HTTP methods directly into the API as function names means any typographical error with the method name will be caught and identified early, for example:

```
>>> requests.foo('http://www.python.org')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'module' object has no attribute 'foo'
>>>
>>> requests.request('foo', 'http://www.python.org')
<Response [403]>
```

The Request and PreparedRequest objects (we're all responsible users)

`__init__.py` exposes `Request`, `PreparedRequest`, and `Response` from `models.py` as part of the main API. Why does `models.Request` even exist? There's already a `url lib.requests.Request` in the standard library, and in `cookies.py` there is specifically a `MockRequest` object that wraps `models.Request` so that it works like `url lib.requests.Request` for `http.cookiejar`.¹⁴ That means whatever methods are needed for the `Request` object to interface with the cookies library are intentionally excluded from `requests.Request`. What is the point of all this extra work?

The extra methods in `MockRequest` (which exists to emulate `url lib.request.Request` for the cookie library) are used by the cookie library to manage cookies. Except for the `get_type()` function (which usually returns “http” or “https” when using `Requests`) and the `unverifiable` property (`True` for our case), they're all related to the URL or the request headers:

Related to the header

`add_unredirected_header()`

Add a new key, value pair to the header.

`get_header()`

Get a specific name in the header dictionary.

`get_new_headers()`

Get the dictionary containing new headers (added by `cookie lib`).

`has_header()`

Check whether a name exists in the header dictionary.

Related to the URL

`get_full_url()`

Does just what it says.

`host and origin_req_host`

Properties that are set by calling the methods `get_host()` and `get_origin_req_host()`, respectively.

`get_host()`

Extract the host from the URL (e.g., `www.python.org` from `https://www.python.org/dev/peps/pep-0008/`).

¹⁴ The module `http.cookiejar` was previously `cookielib` in Python 2, and `urllib.requests.Request` was previously `urllib2.Request` in Python 2.

`get_origin_req_host()`

Call `get_host()`.¹⁵

They're all access functions, except for `MockRequest.add_unredirected_header()`. The `MockRequest` docstring notes that "the original request object is read-only."

In `requests.Request`, data attributes are instead directly exposed. This makes all of the accessor functions unnecessary: to get or set the headers, access `request-instance.headers`. It's just a dictionary. Likewise, the user can just get or change the string URL: `request-instance.url`.

The `PreparedRequest` object is initialized empty, and is populated with a call to `prepared-request-instance.prepare()`, filled with the relevant data (usually from the calling the `Request` object). It's at this point that things like correct capitalization and encoding are applied. The object's contents will, once prepared, be ready to send to the server, but every attribute is still directly exposed. Even `PreparedRequest._cookies` is exposed, although its prepended underscore is a gentle reminder that the attribute is not intended for use outside of the class, without forbidding such access (we are all responsible users).

This choice exposes the objects to user modification, but they are much more readable, and a little bit of extra work inside of `PreparedRequest` corrects capitalization issues and allows use of a dictionary in place of a `CookieJar` (look for the `if isinstance()/else` statement):

```
#
# ... from models.py ...
#

class PreparedRequest():
    #
    # ... skip everything else ...
    #

    def prepare_cookies(self, cookies):
        """Prepares the given HTTP cookie data.

        This function eventually generates a ``Cookie`` header from the
        given cookies using cookiecutter. Due to cookiecutter's design, the header
        will not be regenerated if it already exists, meaning this function
        can only be called once for the life of the
        :class: `PreparedRequest` object. Any subsequent calls
        to ``prepare_cookies`` will have no actual effect, unless the "Cookie"
        header is removed beforehand."""
```

¹⁵ This method makes it possible to handle cross-origin requests (like getting a JavaScript library hosted on a third-party site). It is supposed to return the origin host of the request, defined in [IETF RFC 2965](#).

```

if isinstance(cookies, cookielib.CookieJar):
    self._cookies = cookies
else:
    self._cookies = cookiejar_from_dict(cookies)

cookie_header = get_cookie_header(self._cookies, self)
if cookie_header is not None:
    self.headers['Cookie'] = cookie_header

```

These things may not seem like a big deal, but it's small choices like these that make an API intuitive to use.

Style Examples from Requests

The style examples from Requests are a good example use for sets (which we think aren't used often enough!) and a look at the `requests.status_codes` module, which exists to make the style of the rest of the code simpler by avoiding hardcoded HTTP status codes everywhere else in the code.

Sets and set arithmetic (a nice, Pythonic idiom)

We haven't yet shown an example use of Python sets in action. Python sets behave like sets in math—you can do subtraction, unions (the *or* operator), and intersections (the *and* operator):

```

>>> s1 = set((7,6))
>>> s2 = set((8,7))
>>> s1
{6, 7}
>>> s2
{8, 7}
>>> s1 - s2 # set subtraction
{6}
>>> s1 | s2 # set union
{8, 6, 7}
>>> s1 & s2 # set intersection
{7}

```

Here's one, down toward the end of this function from *cookies.py* (with the label ❷):

```

#
# ... from cookies.py ...
#

def create_cookie(name, value, **kwargs): ❶
    """Make a cookie from underspecified parameters.

    By default, the pair of `name` and `value` will be set for the domain ''
    and sent on every request (this is sometimes called a "supercookie").
    """

```

```

result = dict(
    version=0,
    name=name,
    value=value,
    port=None,
    domain='',
    path='/',
    secure=False,
    expires=None,
    discard=True,
    comment=None,
    comment_url=None,
    rest={'HttpOnly': None},
    rfc2109=False,)

badargs = set(kwargs) - set(result) ❷
if badargs:
    err = 'create_cookie() got unexpected keyword arguments: %s' ❸
    raise TypeError(err % list(badargs)) ❸

result.update(kwargs) ❹
result['port_specified'] = bool(result['port']) ❺
result['domain_specified'] = bool(result['domain'])
result['domain_initial_dot'] = result['domain'].startswith('.')
result['path_specified'] = bool(result['path'])

return cookielib.Cookie(**result) ❻

```

- ❶ The `**kwargs` specification allows the user to provide any or none of the key-word options for a cookie.
- ❷ Set arithmetic! Pythonic. Simple. And in the standard library. On a dictionary, `set()` forms a set of the keys.
- ❸ This is a great example of splitting a long line into two shorter lines that make much better sense. No harm done from the extra `err` variable.
- ❹ The `result.update(kwargs)` updates the `result` dictionary with the key/value pairs in the `kwargs` dictionary, replacing existing pairs or creating ones that didn't exist.
- ❺ Here the call to `bool()` coerces the value to `True` if the object is truthy (meaning it evaluates to `True`—in this case, `bool(result['port'])` evaluates to `True` if it's not `None` and it's not an empty container).
- ❻ The signature to initialize `cookielib.Cookie` is actually 18 positional arguments and one keyword argument (`rfc2109` defaults to `False`). It's impossible for us average humans to memorize which position has which value, so here Requests

takes advantage of being able to assign positional arguments by name as keyword arguments, sending the whole dictionary.

Status codes (readability counts)

The entire `status_codes.py` exists to create an object that can look up status codes by attribute. We're showing the definition of the lookup dictionary in `status_codes.py` first, and then an excerpt of code that uses it from `sessions.py`:

```
#
# ... excerpted from requests/status_codes.py ...
#

_codes = {

    # Informational.
    100: ('continue',),
    101: ('switching_protocols',),
    102: ('processing',),
    103: ('checkpoint',),
    122: ('uri_too_long', 'request_uri_too_long'),
    200: ('ok', 'okay', 'all_ok', 'all_okay', 'all_good', '\\o/', '✓'), ❶
    201: ('created',),
    202: ('accepted',),
    #
    # ... skipping ...
    #

    # Redirection.
    300: ('multiple_choices',),
    301: ('moved_permanently', 'moved', '\\o-'),
    302: ('found',),
    303: ('see_other', 'other'),
    304: ('not_modified',),
    305: ('use_proxy',),
    306: ('switch_proxy',),
    307: ('temporary_redirect', 'temporary_moved', 'temporary'),
    308: ('permanent_redirect',
        'resume_incomplete', 'resume'), # These 2 to be removed in 3.0 ❷

    #
    # ... skipping the rest ...
    #
}

codes = LookupDict(name='status_codes') ❸

for code, titles in _codes.items():
    for title in titles:
        setattr(codes, title, code) ❹
        if not title.startswith('\\o'):
            setattr(codes, title.upper(), code) ❺
```

- ❶ All of these options for an OK status will become keys in the lookup dictionary. Except for the happy person (\\o/) and the check mark (✓).
- ❷ The deprecated values are on a separate line so that the future delete will be clean and obvious in version control.
- ❸ The LookupDict allows dot-access of its elements like in the next line.
- ❹ `codes.ok == 200` and `codes.okay == 200`.
- ❺ And also `codes.OK == 200` and `codes.OKAY == 200`.

All of this work for the status codes was to make the lookup dictionary `codes`. Why? Instead of very typo-prone hardcoded integers all over the code, this is easy to read, with all of the code numbers localized in a single file. Because it starts as a dictionary keyed on the status codes, each status code integer only exists once. The possibility of typos is much, much lower than if this were just a bunch of global variables manually embedded into a namespace.

And converting the keys into attributes instead of using them as strings in a dictionary again reduces the risk of typographical errors. Here's the example in *sessions.py* that's so much easier to read with words than numbers:

```
#
# ... from sessions.py ...
#     Truncated to only show relevant content.
#
from .status_codes import codes ❶

class SessionRedirectMixin(object): ❷
    def resolve_redirects(self, resp, req, stream=False, timeout=None,
                        verify=True, cert=None, proxies=None,
                        **adapter_kwargs):
        """Receives a Response. Returns a generator of Responses."""

        i = 0
        hist = [] # keep track of history

        while resp.is_redirect: ❸
            prepared_request = req.copy()

            if i > 0:
                # Update history and keep track of redirects.
                hist.append(resp)
                new_hist = list(hist)
                resp.history = new_hist

            try:
```



```

        resp.content # Consume socket so it can be released
    except (ChunkedEncodingError, ContentDecodingError, RuntimeError):
        resp.raw.read(decode_content=False)

    if i >= self.max_redirects:
        raise TooManyRedirects(
            'Exceeded %s redirects.' % self.max_redirects
        )

    # Release the connection back into the pool.
    resp.close()

    #
    # ... skipping content ...
    #

    # http://tools.ietf.org/html/rfc7231#section-6.4.4
    if (resp.status_code == codes.see_other and ❷
        method != 'HEAD'):
        method = 'GET'

    # Do what the browsers do, despite standards...
    # First, turn 302s into GETs.
    if resp.status_code == codes.found and method != 'HEAD': ❸
        method = 'GET'

    # Second, if a POST is responded to with a 301, turn it into a GET.
    # This bizarre behavior is explained in Issue 1704.
    if resp.status_code == codes.moved and method == 'POST': ❹
        method = 'GET'

    #
    # ... etc. ...
    #

```

- ❶ Here's where the status code lookup codes is imported.
- ❷ We describe mixin classes later in “Mixins (also one honking great idea)” on page 153. This mixin provides redirect methods for the main Session class, which is defined in this same file but not shown in our excerpt.
- ❸ We're entering a loop that's following the redirects for us to get at the content we want. The entire loop logic is deleted from this excerpt for brevity.
- ❹ Status codes as text are so much more readable than unmemorable integers: `codes.see_other` would otherwise be 303 here.
- ❺ And `codes.found` would be 302, and `codes.moved` would be 301. So the code is self-documenting; we can tell meaning from the variable names; and we've avoided

ded the possibility of littering the code with typographical errors by using dot-notation instead of a dictionary to look up strings (e.g., `codes.found` instead of `codes["found"]`).

Werkzeug

To read Werkzeug, we need to know a little about how web servers communicate with applications. The next paragraphs try to give as short an overview as possible.

Python’s interface for web application-to-server interaction, WSGI, is defined in PEP 333, which was written by Phillip J. Eby in 2003.¹⁶ It specifies how a web server (like Apache) communicates with a Python application or framework:

1. The server will call the application once per HTTP request (e.g., “GET” or “POST”) it receives.
2. That application will return an iterable of bytestrings that the server will use to respond to the HTTP request.
3. The specification also says the application will take two parameters—for example, `webapp(environ, start_response)`. The `environ` parameter will contain all of the data associated with the request, and the `start_response` parameter will be a function or other callable object that will be used to send back header (e.g., `('Content-type', 'text/plain')`) and status (e.g., `200 OK`) information to the server.

That summary glosses over about a half-dozen pages of additional detail. In the middle of PEP 333 is this aspirational statement about the new standard making modular web frameworks possible:

If middleware can be both simple and robust, and WSGI is widely available in servers and frameworks, it allows for the possibility of an entirely new kind of Python web application framework: one consisting of loosely-coupled WSGI middleware components. Indeed, existing framework authors may even choose to refactor their frameworks’ existing services to be provided in this way, becoming more like libraries used with WSGI, and less like monolithic frameworks. This would then allow application developers to choose “best-of-breed” components for specific functionality, rather than having to commit to all the pros and cons of a single framework.

Of course, as of this writing, that day is doubtless quite far off. In the meantime, it is a sufficient short-term goal for WSGI to enable the use of any framework with any server.

¹⁶ Since then, PEP 333 has been superseded by a specification updated to include some Python 3-specific details, [PEP 3333](#). For a digestible but very thorough introduction, we recommend Ian Bicking’s [WSGI tutorial](#).

About four years after, in 2007, Armin Ronacher released Werkzeug, with the intent of filling that hopeful need for a WSGI library that can be used to make WSGI applications and middleware components.

Werkzeug is the largest package we are reading, so we'll highlight just a few of its design choices.

Reading Code in a Toolkit

A software toolkit is a collection of compatible utilities. In Werkzeug's case, they're all related to WSGI applications. A good way to understand the distinct utilities and what they're for is to look at the unit tests, and that's how we'll approach reading Werkzeug's code.

Get Werkzeug from GitHub:

```
$ git clone https://github.com/pallets/werkzeug.git
$ virtualenv -p python3 venv
$ source venv/bin/activate
(venv)$ cd werkzeug
(venv)$ pip install --editable .
(venv)$ py.test tests # Run the unit tests
```

Read Werkzeug's documentation

[Werkzeug's documentation](#) lists the main things it provides—an implementation of the WSGI 1.0 ([PEP 333](#)) specification, a URL routing system, the capability to parse and dump HTTP headers, objects that represent HTTP requests and HTTP responses, session and cookie support, file uploads, and other utilities and community additions. Plus, a full-featured debugger.

The tutorials are good, but we're using the API documentation instead to see more of the library's components. The next section takes from Werkzeug's [wrappers](#) and [routing documentation](#).

Use Werkzeug

Werkzeug provides utilities for WSGI applications, so to learn what Werkzeug provides, we can start with a WSGI application, and then use a few of Werkzeug's utilities. This first application is a slightly changed version of what's in PEP 333, and doesn't use Werkzeug yet. The second one does the same thing as the first, but using Werkzeug:

```
def wsgi_app(environ, start_response):
    headers = [('Content-type', 'text/plain'), ('charset', 'utf-8')]
    start_response('200 OK', headers)
    yield 'Hello world.'
```

```
# This app does the same thing as the one above:
response_app = werkzeug.Response('Hello world!')
```

Werkzeug implements a `werkzeug.Client` class to stand in for a real web sever when doing one-off testing like this. The client response will have the type of the `response_wrapper` argument. In this code, we create clients and use them to call the WSGI applications we made earlier. First, the plain WSGI app (but with the response parsed into a `werkzeug.Response`):

```
>>> import werkzeug
>>> client = werkzeug.Client(wsgi_app, response_wrapper=werkzeug.Response)
>>> resp=client.get("?answer=42")
>>> type(resp)
<class 'werkzeug.wrappers.Response'>
>>> resp.status
'200 OK'
>>> resp.content_type
'text/plain'
>>> print(resp.data.decode())
Hello world.
```

Next, using the `werkzeug.Response` WSGI app:

```
>>> client = werkzeug.Client(response_app, response_wrapper=werkzeug.Response)
>>> resp=client.get("?answer=42")
>>> print(resp.data.decode())
Hello world!
```

The `werkzeug.Request` class provides the contents of the environment dictionary (the `environ` argument above to `wsgi_app()`) in a form that's easier to use. It also provides a decorator to convert a function that takes a `werkzeug.Request` and returns a `werkzeug.Response` into a WSGI app:

```
>>> @werkzeug.Request.application
... def wsgi_app_using_request(request):
...     msg = "A WSGI app with:\n method: {}\n path: {}\n query: {}\n"
...     return werkzeug.Response(
...         msg.format(request.method, request.path, request.query_string))
... 
```

which, when used, gives:

```
>>> client = werkzeug.Client(
...     wsgi_app_using_request, response_wrapper=werkzeug.Response)
>>> resp=client.get("?answer=42")
>>> print(resp.data.decode())
A WSGI app with:
method: GET
path: /
query: b'answer=42'
```

So, we know how to use the `werkzeug.Request` and `werkzeug.Response` objects. The other thing that was featured in the documentation was the routing. Here's an excerpt that uses it—callout numbers identify both the pattern and its match:

```
>>> import werkzeug
>>> from werkzeug.routing import Map, Rule
>>>
>>> url_map = Map([ ❶
...     Rule('/', endpoint='index'), ❷
...     Rule('/<any("Robin","Galahad","Arthur"):person>', endpoint='ask'), ❸
...     Rule('/<other>', endpoint='other') ❹
... ])

>>> env = werkzeug.create_environ(path='/shouldnt/match') ❺
>>> urls = url_map.bind_to_environ(env)
>>> urls.match()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "[...path...]/werkzeug/werkzeug/routing.py", line 1569, in match
    raise NotFound()
werkzeug.exceptions.NotFound: 404: Not Found
```

- ❶ The `werkzeug.Routing.Map` provides the main routing functions. The rule matching is done in order; the first rule to match is the one selected.
- ❷ When there are no angle-bracket terms in the rule's placeholder string, it only matches on an exact match, and the second result from `urls.match()` is an empty dictionary:

```
>>> env = werkzeug.create_environ(path='/')
>>> urls = url_map.bind_to_environ(env)
>>> urls.match()
('index', {})
```

- ❸ Otherwise, the second entry is a dictionary mapping the named terms in the rule to their value—for example, mapping 'person' to the value 'Galahad':

```
>>> env = werkzeug.create_environ(path='/Galahad?favorite+color')
>>> urls = url_map.bind_to_environ(env)
>>> urls.match()
('ask', {'person': 'Galahad'})
```

- ❹ Note that 'Galahad' could have matched the route named 'other', but it did not—but 'Lancelot' did—because the first rule to match the pattern is chosen:

```
>>> env = werkzeug.create_environ(path='/Lancelot')
>>> urls = url_map.bind_to_environ(env)
>>> urls.match()
('other', {'other': 'Lancelot'})
```

- ⑤ And an exception is raised if there are no matches at all in the list of rules:

```
>>> env = werkzeug.test.create_environ(path='/shouldnt/match')
>>> urls = url_map.bind_to_environ(env)
>>> urls.match()
Traceback (most recent call last):
File "<stdin>", line 1, in <module>
File "[...path...]/werkzeug/werkzeug/routing.py", line 1569, in match
raise NotFound()
werkzeug.exceptions.NotFound: 404: Not Found
```

You'd use the map to route a request to an appropriate endpoint. The following code continues on from the preceding example to do this:

```
@werkzeug.Request.application
def send_to_endpoint(request):
    urls = url_map.bind_to_environ(request)
    try:
        endpoint, kwargs = urls.match()
        if endpoint == 'index':
            response = werkzeug.Response("You got the index.")
        elif endpoint == 'ask':
            questions = dict(
                Galahad='What is your favorite color?',
                Robin='What is the capital of Assyria?',
                Arthur='What is the air-speed velocity of an unladen swallow?')
            response = werkzeug.Response(questions[kwargs['person']])
        else:
            response = werkzeug.Response("Other: {other}".format(**kwargs))
    except (KeyboardInterrupt, SystemExit):
        raise
    except:
        response = werkzeug.Response(
            'You may not have gone where you intended to go,\n'
            'but I think you have ended up where you needed to be.',
            status=404
        )
    return response
```

To test it, use the `werkzeug.Client` again:

```
>>> client = werkzeug.Client(send_to_endpoint, response_wrapper=werkzeug.Response)
>>> print(client.get("/").data.decode())
You got the index.
>>>
>>> print(client.get("Arthur").data.decode())
What is the air-speed velocity of an unladen swallow?
>>>
>>> print(client.get("42").data.decode())
Other: 42
>>>
>>> print(client.get("time/lunchtime").data.decode()) # no match
```

You may **not** have gone where you intended to go,
but I think you have ended up where you needed to be.

Read Werkzeug's code

When test coverage is good, you can learn what a library does by looking at the unit tests. The caveat is that with unit tests, you're intentionally looking at the “trees” and not the “forest”—exploring obscure use cases intended to ensure the code doesn't break, rather than looking for interconnections between modules. This should be OK for a toolkit like Werkzeug, which we expect to have modular, loosely coupled components.

Because we familiarized ourselves with how the routing and the request and response wrappers work, *werkzeug/test_routing.py* and *werkzeug/test_wrappers.py* are good choices to read for now.

When we first open *werkzeug/test_routing.py*, we can quickly look for interconnection between the modules by searching through the entire file for the imported objects. Here are all of the `import` statements:

```
import pytest ❶

import uuid ❷

from tests import strict_eq ❸

from werkzeug import routing as r ❹
from werkzeug.wrappers import Response ❺
from werkzeug.datastructures import ImmutableDict, MultiDict ❻
from werkzeug.test import create_environ ❼
```

- ❶ Of course, `pytest` is used here for testing.
- ❷ The `uuid` module is used in just one function, `test_uuid_converter()`, to confirm that the conversion from string to a `uuid.UUID` object (the Universal Unique Identifier string uniquely identifying objects on the Internet) works.
- ❸ The `strict_eq()` function is used often, and defined in *werkzeug/tests/__init__.py*. It's for testing, and is only necessary because in Python 2 there used to be implicit type conversion between Unicode and byte strings, but relying on this breaks things in Python 3.
- ❹ The `werkzeug.routing` module is the one that's being tested.
- ❺ The `Response` object is used in just one function, `test_dispatch()`, to confirm that `werkzeug.routing.MapAdapter.dispatch()` passes the correct information through to the dispatched WSGI application.

- ⑥ These dictionary objects are used only once each, `ImmutableDict` to confirm that an immutable dictionary in `werkzeug.routing.Map` is indeed immutable, and `MultiDict` to provide multiple keyed values to the URL builder and confirm that it still builds the correct URL.
- ⑦ The `create_environ()` function is for testing—it creates a WSGI environment without having to use an actual HTTP request.

The point of doing the quick searching was to quickly see the interconnection between modules. What we found out was that `werkzeug.routing` imports some special data structures, and that's all. The rest of the unit tests show the scope of the routing module. For example, non-ASCII characters can be used:

```
def test_environ_nonascii_pathinfo():
    environ = create_environ(u'/лошадь')
    m = r.Map([
        r.Rule(u'/', endpoint='index'),
        r.Rule(u'/лошадь', endpoint='horse')
    ])
    a = m.bind_to_environ(environ)
    strict_eq(a.match(u'/'), ('index', {}))
    strict_eq(a.match(u'/лошадь'), ('horse', {}))
    pytest.raises(r.NotFound, a.match, u'/барык')
```

There are tests to build and parse URLs, and even utilities to find the closest available match, when there wasn't an actual match. You can even do all kinds of crazy custom processing when handling the type conversions/parsing from the path and the URL string:

```
def test_converter_with_tuples():
    """
    Regression test for https://github.com/pallets/werkzeug/issues/709
    """
    class TwoValueConverter(r.BaseConverter):
        def __init__(self, *args, **kwargs):
            super(TwoValueConverter, self).__init__(*args, **kwargs)
            self.regex = r'(\w\w+)/(\w\w+)'

        def to_python(self, two_values):
            one, two = two_values.split('/')
            return one, two

        def to_url(self, values):
            return "%s/%s" % (values[0], values[1])

    map = r.Map([
        r.Rule('/<two:foo>/', endpoint='handler')
    ], converters={'two': TwoValueConverter})
    a = map.bind('example.org', '/')
```



```
route, kwargs = a.match('/qwertyuiop/')
assert kwargs['foo'] == ('qwerty', 'yuiop')
```

Similarly, *werkzeug/test_wrappers.py* does not import much. Reading through the tests gives an example of the scope of available functionality for the Request object—cookies, encoding, authentication, security, cache timeouts, and even multilanguage encoding:

```
def test_modified_url_encoding():
    class ModifiedRequest(wrappers.Request):
        url_charset = 'euc-kr'

    req = ModifiedRequest.from_values(u'/?foo=정상처리'.encode('euc-kr'))
    strict_eq(req.args['foo'], u'정상처리')
```

In general, reading the tests provides a way to see the details of what the library provides. Once satisfied that we have an idea of what Werkzeug is, we can move on.

Tox in Werkzeug

Tox is a Python command-line tool that uses virtual environments to run tests. You can run it on your own computer (tox on the command line), so long as the Python interpreters you're using are already installed. It's integrated with GitHub, so if you have a *tox.ini* file in the top level of your repository, like Werkzeug does, it will automatically run tests on every commit.

Here is Werkzeug's entire *tox.ini* configuration file:

```
[tox]
envlist = py{26,27,py,33,34,35}-normal, py{26,27,33,34,35}-uwsgi

[testenv]
passenv = LANG
deps=
# General
    pyopenssl
    greenlet
    pytest
    pytest-xprocess
    redis
    requests
    watchdog
    uwsgi: uwsgi

# Python 2
    py26: python-memcached
    py27: python-memcached
    pypy: python-memcached

# Python 3
    py33: python3-memcached
```

```

py34: python3-memcached
py35: python3-memcached

whitelist_externals=
    redis-server
    memcached
    uwsgi

commands=
    normal: py.test []
    uwsgi: uwsgi
           --pyrun {envbindir}/py.test
           --pyargv -kUWSGI --cache2=name=werkzeugtest,items=20 --master

```

Style Examples from Werkzeug

Most of the major style points we made in [Chapter 4](#) have already been covered. The first style example we chose shows an elegant way to guess types from a string, and the second one makes a case for using the `VERBOSE` option when defining long regular expressions—so that other people can tell what the expression does without having to spend time thinking through it.

Elegant way to guess type (if the implementation is easy to explain, it may be a good idea)

If you're like most of us, you've had to parse text files and convert content to various types. This solution is particularly Pythonic, so we wanted to include it:

```

_PYTHON_CONSTANTS = {
    'None':    None,
    'True':    True,
    'False':   False
}

def _pythonize(value):
    if value in _PYTHON_CONSTANTS: ❶
        return _PYTHON_CONSTANTS[value]
    for convert in int, float: ❷
        try: ❸
            return convert(value)
        except ValueError:
            pass
    if value[:1] == value[-1:] and value[0] in '"\': ❹
        value = value[1:-1]
    return text_type(value) ❺

```

- ❶ Key lookup for Python dictionaries uses hash mapping, just like set lookup. Python doesn't have switch case statements. (They were proposed and rejected

for lack of popularity in [PEP 3103](#).) Instead, Python users use `if/elif/else`, or as shown here, the very Pythonic option of a dictionary lookup.

- ❷ Note that the first conversion attempt is to the more restrictive type, `int`, before attempting the conversion to `float`.
- ❸ It is also Pythonic to use `try/except` statements to infer type.
- ❹ This part is necessary because the code is in *werkzeug/routing.py*, and the string being parsed is part of a URL. It's checking for quotes and unquoting the value.
- ❺ `text_type` converts strings to Unicode in a way that's both Python 2 and Python 3 compatible. It's basically the same thing as the `u()` function highlighted in [“HowDoI” on page 93](#).

Regular expressions (readability counts)

If you use lengthy regular expressions in your code, please use the `re.VERBOSE`¹⁷ option and make it comprehensible to the rest of us humans, like this snippet from *werkzeug/routing.py*:

```
import re

_rule_re = re.compile(r'''
    (?P<static>[^<]*)           # static rule data
    <
    (? :
        (?P<converter>[a-zA-Z_][a-zA-Z0-9_]*) # converter name
        (? : \((?P<args>.*?)\))?           # converter arguments
        \:                                 # variable delimiter
    )?
    (?P<variable>[a-zA-Z_][a-zA-Z0-9_]*)    # variable name
    >
''', re.VERBOSE)
```

Structure Examples from Werkzeug

The first two examples related to structure demonstrate Pythonic ways to leverage dynamic typing. We cautioned against reassigning a variable to different values in [“Dynamic Typing” on page 68](#) but didn't mention any benefits. One of them is the ability to use any type of object that behaves in the expected way—*duck typing*. Duck

¹⁷ `re.VERBOSE` allows you to write more readable regular expressions by changing the way whitespace is treated, and by allowing comments. Read more in the [re documentation](#).

typing approaches types with the philosophy: “If it looks like a duck¹⁸ and quacks like a duck, then it’s a duck”.

They both play in different ways on ways that objects can be callable without being functions: `cached_property.__init__()` allows initialization of a class instance to be used like an ordinary function call, and `Response.__call__()` allows a `Response` instance to itself be called like a function.

The last excerpt uses Werkzeug’s implementation of some mixin classes (that each define a subset of the functionality in Werkzeug’s `Request` object) to discuss why they’re a honking great idea.

Class-based decorators (a Pythonic use of dynamic typing)

Werkzeug makes use of duck typing to make the `@cached_property` decorator. When we talked about `property` when describing the `Tablib` project, we talked about it like it’s a function. Usually decorators *are* functions, but because there is no enforcement of type, they can be any callable: `property` is actually a class. (You can tell it’s intended to be used like a function because it is not capitalized, like PEP 8 says class names should be.) When written like a function call (`property()`), `property.__init__()` will be called to initialize and return a `property` instance—a class, with an appropriately defined `__init__()` method, works as a callable. Quack.

The following excerpt contains the entire definition of `cached_property`, which subclasses the `property` class. The documentation within `cached_property` speaks for itself. When it is used to decorate `BaseRequest.form` in the code we just saw, the *instance*.`form` will have the type `cached_property` and will behave like a dictionary as far as the user is concerned, because both the `__get__()` and `__set__()` methods are defined. The first time `BaseRequest.form` is accessed, it will read its form data (if it exists) once, and then store the data in *instance*.`form.__dict__` to be accessed in the future:

```
class cached_property(property):

    """A decorator that converts a function into a lazy property. The
    function wrapped is called the first time to retrieve the result,
    and then that calculated result is used the next time you access
    the value::

        class Foo(object):

            @cached_property
            def foo(self):
                # calculate something important here
```

¹⁸ That is, if it’s callable, or iterable, or has the correct method defined ...

```
    return 42
```

The class has to have a `__dict__` in order for this property to work.

```
"""
```

```
# implementation detail: A subclass of Python's built-in property  
# decorator, we override __get__ to check for a cached value. If one  
# choses to invoke __get__ by hand, the property will still work as  
# expected because the lookup logic is replicated in __get__ for  
# manual invocation.
```

```
def __init__(self, func, name=None, doc=None):  
    self.__name__ = name or func.__name__  
    self.__module__ = func.__module__  
    self.__doc__ = doc or func.__doc__  
    self.func = func  
  
def __set__(self, obj, value):  
    obj.__dict__[self.__name__] = value  
  
def __get__(self, obj, type=None):  
    if obj is None:  
        return self  
    value = obj.__dict__.get(self.__name__, _missing)  
    if value is _missing:  
        value = self.func(obj)  
        obj.__dict__[self.__name__] = value  
    return value
```

Here it is in action:

```
>>> from werkzeug.utils import cached_property  
>>>  
>>> class Foo(object):  
...     @cached_property  
...     def foo(self):  
...         print("You have just called Foo.foo()!")  
...         return 42  
...  
>>> bar = Foo()  
>>>  
>>> bar.foo  
You have just called Foo.foo()!  
42  
>>> bar.foo  
42  
>>> bar.foo # Notice it doesn't print again...  
42
```

Response.__call__

The Response class is built using features mixed into the BaseResponse class, just like Request. We will highlight its user interface and won't show the actual code, just the docstring for BaseResponse, to show the usage details:

```
class BaseResponse(object):

    """Base response class. The most important fact about a response object
    is that it's a regular WSGI application. It's initialized with a couple
    of response parameters (headers, body, status code, etc.) and will start a
    valid WSGI response when called with the environ and start response
    callable.

    Because it's a WSGI application itself, processing usually ends before the
    actual response is sent to the server. This helps debugging systems
    because they can catch all the exceptions before responses are started.

    Here is a small example WSGI application that takes advantage of the
    response objects::

        from werkzeug.wrappers import BaseResponse as Response

        def index(): ❶
            return Response('Index page')

        def application(environ, start_response): ❷
            path = environ.get('PATH_INFO') or '/'
            if path == '/':
                response = index() ❸
            else:
                response = Response('Not Found', status=404) ❹
            return response(environ, start_response) ❺
    """
    # ... etc. ...
```

- ❶ In the example from the docstring, index() is the function that will be called in response to the HTTP request. The response will be the string “Index page”.
- ❷ This is the signature required for a WSGI application, as specified in PEP 333/ PEP 3333.
- ❸ Response subclasses BaseResponse, so response is an instance of BaseResponse.
- ❹ See how the 404 response just requires the status keyword to be set.
- ❺ And, voilà, the response instance is itself callable, with all of the accompanying headers and details set to sensible default values (or overrides in the case where the path is not “/”).

So, how is an instance of a class callable? Because the `BaseRequest.__call__` method has been defined. We show just that method in the following code example.

```
class BaseResponse(object):
    #
    # ... skip everything else ...
    #

    def __call__(self, environ, start_response): ❶
        """Process this response as WSGI application.

        :param environ: the WSGI environment.
        :param start_response: the response callable provided by the WSGI
                               server.
        :return: an application iterator
        """
        app_iter, status, headers = self.get_wsgi_response(environ)
        start_response(status, headers) ❷
        return app_iter ❸
```

- ❶ Here's the signature to make `BaseResponse` instances callable.
- ❷ Here's where the `start_response` function call requirement of WSGI apps is satisfied.
- ❸ And here's where the iterable of bytes is returned.

The lesson here is this: if it's possible in the language, why not do it? We promise after realizing we could add a `__call__()` method to any object and make it callable, we were inspired to go back to the original documentation for a good re-read of [Python's data model](#).

Mixins (also one honking great idea)

Mixins in Python are classes that are intended to add a specific functionality—a handful of related attributes—to a class. Python, unlike Java, allows for multiple inheritance. This means that the following paradigm, where a half-dozen different classes are subclassed simultaneously, is a possible way to modularize different functionality into separate classes. “Namespaces,” sort of.

Modularization like this is useful in a utility library like Werkzeug because it communicates to the user which functions are related and not related: the developer can be confident that attributes in one mixin are not going to be modified by any functions in another mixin.



In Python, there isn't anything special to identify a mixin, other than the convention of appending `Mixin` to the end of the class name. This means if you don't want to pay attention to the order of method resolution, all of the mixins' methods should have distinct names.

In Werkzeug, sometimes methods in a mixin may require certain attributes to be present. These requirements are usually documented in the mixin's docstring:

```
# ... in werkzeug/wrappers.py
```

```
class UserAgentMixin(object): ❶
```

```
    """Adds a `user_agent` attribute to the request object which contains
    the parsed user agent of the browser that triggered the request as a
    :class:`~werkzeug.useragents.UserAgent` object.
    """
```

```
@cached_property
```

```
def user_agent(self):
```

```
    """The current user agent."""
```

```
    from werkzeug.useragents import UserAgent
```

```
    return UserAgent(self.environ) ❷
```

```
class Request(BaseRequest, AcceptMixin, ETagRequestMixin,
              UserAgentMixin, AuthorizationMixin, ❸
              CommonRequestDescriptorsMixin):
```

```
    """Full featured request object implementing the following mixins:
```

```
    - :class:`AcceptMixin` for accept header parsing
    - :class:`ETagRequestMixin` for etag and cache control handling
    - :class:`UserAgentMixin` for user agent introspection
    - :class:`AuthorizationMixin` for http auth handling
    - :class:`CommonRequestDescriptorsMixin` for common headers
    """
```

❹

- ❶ There is nothing special about the `UserAgentMixin`; it subclasses `object`, though, which is the default in Python 3, highly recommended for compatibility in Python 2, and which should be done explicitly because, well, “explicit is better than implicit.”
- ❷ `UserAgentMixin.user_agent` assumes there is a `self.environ` attribute.
- ❸ When included in the list of base classes for `Request`, the attribute it provides becomes accessible via `Request(environ).user_agent`.

- ④ Nothing else—this is the entire body of the definition of `Request`. All functionality is provided by the base class or the mixins. Modular, pluggable, and as froody as Ford Prefect.

New-Style Classes and object

The base class `object` adds default attributes that other built-in options rely on. Classes that don't inherit from `object` are called “old-style classes” or “classic classes” and are removed in Python 3. It's the default to inherit from `object` in Python 3, meaning all Python 3 classes are “new-style classes.” New-style classes are available in Python 2.7 (actually with their current behavior since Python 2.3), but the inheritance must be written explicitly, and (we think) should always be written.

There are more details in the Python documentation on [new-style classes](#), a tutorial [here](#), and a technical history of their creation in [this post](#). Here are some of the differences (in Python 2.7; all classes are new-style in Python 3):

```
>>> class A(object):
...     """New-style class, subclassing object."""
...
>>> class B:
...     """Old-style class."""
...
>>> dir(A)
['_class_', '__delattr__', '__dict__', '__doc__', '__format__',
 '__getattr__', '__hash__', '__init__', '__module__', '__new__',
 '__reduce__', '__reduce_ex__', '__repr__', '__setattr__', '__sizeof__',
 '__str__', '__subclasshook__', '__weakref__']
>>>
>>> dir(B)
['__doc__', '__module__']
>>>
>>> type(A)
<type 'type'>
>>> type(B)
<type 'classobj'>
>>>
>>> import sys
>>> sys.getsizeof(A()) # The size is in bytes.
64
>>> sys.getsizeof(B())
72
```

Flask

Flask is a web microframework that combines Werkzeug and Jinja2, both by Armin Ronacher. It was created as a joke and released on April Fool's Day, 2010, but quickly

became one of Python’s most popular web frameworks. He had released Werkzeug a few years earlier in 2007 as a “Swiss Army knife of Python web development,” and (we presume) was probably a little frustrated at its slow adoption. The idea for Werkzeug was to decouple the WSGI from everything else so that developers could plug in their own choice of utilities. Little did he know how much we’d appreciate a few more “rails.”¹⁹

Reading Code in a Framework

A software framework is just like a physical framework—it provides the underlying structure to build a WSGI²⁰ application: the library user provides components that the main Flask application will run. Our goal in reading will be to understand the framework structure and what precisely it provides.

Get Flask from GitHub:

```
$ git clone https://github.com/pallets/flask.git
$ virtualenv venv # Python 3 is usable but discouraged for now
$ source venv/bin/activate
(venv)$ cd flask
(venv)$ pip install --editable .
(venv)$ pip install -r test-requirements.txt # Required for unit tests
(venv)$ py.test tests # Run the unit tests
```

Read Flask’s documentation

Flask’s [online documentation](#) starts out with a seven-line implementation of a web app, and then summarizes Flask: it’s a Unicode-based WSGI-compliant framework that uses Jinja2 for HTML templating and Werkzeug for WSGI utilities like URL routing. It also has built-in tools for development and testing. There are also tutorials, so the next step is easy.

Use Flask

We can run the flaskr example that we downloaded with the GitHub repository. The documents say it’s a small blog site. From within the top *flask* directory:

```
(venv)$ cd examples/flaskr/
(venv)$ py.test test_flaskr.py # Tests should pass
(venv)$ export FLASK_APP=flaskr
```

19 A reference to Ruby on Rails, which popularized web frameworks, and which is much more similar to Django’s style of “everything included,” rather than Flask’s style of “nearly nothing included” (until you add plugins). Django is a great choice when the things you want included are the things Django provides—it was made, and is fantastic for, hosting an online newspaper.

20 WSGI is a Python standard, defined in [PEP 333](#) and [PEP 3333](#) for how an application can communicate with a web server.

```
(venv)$ flask initdb
(venv)$ flask run
```

Read Flask's code

The ultimate goal of Flask is to create a web application, so really it isn't so different from the command-line applications Diamond and HowDoI. Rather than another diagram tracing the flow of function calls through the code, this time we'll step through Flask by running the flaskr example app with a debugger; we'll use `pdb`—the Python debugger—in the standard library.

First, add a breakpoint to `flaskr.py`, which will be activated when that point in the code is reached, causing the interactive session to enter the debugger:

```
@app.route('/')
def show_entries():
    import pdb; pdb.set_trace() ## This line is the breakpoint.
    db = get_db()
    cur = db.execute('select title, text from entries order by id desc')
    entries = cur.fetchall()
    return render_template('show_entries.html', entries=entries)
```

Next, close the file and type `python` on the command line to enter an interactive session. Rather than starting a server, use Flask's internal testing utilities to simulate an HTTP GET request to the `/` location where we just placed the debugger:

```
>>> import flaskr
>>> client = flaskr.app.test_client()
>>> client.get('/')
> /[...] truncated path ...]/flask/examples/flaskr/flaskr.py(74)show_entries()
-> db = get_db()
(Pdb)
```

The last three lines are from `pdb`: we see the path (to `flaskr.py`), the line number (74), and the method name (`show_entries()`) where we stopped. The line (`-> db = get_db()`) shows the statement that will be executed next if we were to step forward in the debugger. And the `(Pdb)` prompt reminds us that we are using the `pdb` debugger.

We can navigate up or down the stack²¹ by typing `u` or `d`, respectively, at the command prompt. See the [pdb documentation](#) under the header “Debugger Commands” for a complete list of the commands you can type. We can also type variable names to see

²¹ The Python call stack contains the instructions that are in progress, being run by the Python interpreter. So if function `f()` calls function `g()`, then function `f()` will go on the stack first, and `g()` will be pushed on top of `f()` when it's called. When `g()` returns, it is *popped* off of (removed from) the stack, and then `f()` will continue where it left off. It is called a stack, because conceptually it works the same way a dish washer will approach a stack of plates—new ones go on the top, and you always deal with the top ones first.

them, and any other Python command; we can even set the variables to different values before we continue on in the code.

If we go up the stack one step, we see what called the `show_entries()` function (with the breakpoint we just installed): it's a `flask.app.Flask` object with a lookup dictionary named `view_functions` that maps string names (like `'show_entries'`) to functions. We also see the `show_entries()` function was called with `**req.view_args`. We can check what `req.view_args` is from the interactive debugger command line by just typing its name (it's the empty dictionary — `{}`, meaning no arguments):

```
(Pdb) u
> /[ ... truncated path ...]/flask/flask/app.py(1610)dispatch_request()
-> return self.view_functions[rule.endpoint](**req.view_args)
(Pdb) type(self)
<class 'flask.app.Flask'>
(Pdb) type(self.view_functions)
<type 'dict'>
(Pdb) self.view_functions
{'add_entry': <function add_entry at 0x108198230>,
 'show_entries': <function show_entries at 0x1081981b8>, [... truncated ...]
 'login': <function login at 0x1081982a8>}
(Pdb) rule.endpoint
'show_entries'
(Pdb) req.view_args
{}
```

We can simultaneously follow along through the source code, if we want to, by opening the appropriate file and going to the stated line. If we keep going up the stack, we can see where the WSGI application is called:

```
(Pdb) u
> /[ ... truncated path ...]/flask/flask/app.py(1624)full_dispatch_request()
-> rv = self.dispatch_request()
(Pdb) u
> /[ ... truncated path ...]/flask/flask/app.py(1973)wsgi_app()
-> response = self.full_dispatch_request()
(Pdb) u
> /[ ... truncated path ...]/flask/flask/app.py(1985).__call__()
-> return self.wsgi_app(environ, start_response)
```

If we type `u` any more, we end up in the testing module, which was used to create the fake client without having to start a server—we've gone as far up the stack as we want to go. We learned that the application `flaskr` is dispatched from within an instance of the `flask.app.Flask` class, on line 1985 of `flask/flask/app.py`. Here is the function:

```
class Flask:
    ## ~~ ... skip lots of definitions ...

    def wsgi_app(self, environ, start_response):
        """The actual WSGI application. ... skip other documentation ...
        """
```

```

    ctx = self.request_context(environ)
    ctx.push()
    error = None
    try:
        try:
            response = self.full_dispatch_request() ❶
        except Exception as e:
            error = e
            response = self.make_response(self.handle_exception(e))
        return response(environ, start_response)
    finally:
        if self.should_ignore_error(error):
            error = None
        ctx.auto_pop(error)

def __call__(self, environ, start_response):
    """Shortcut for :attr:`wsgi_app`. """
    return self.wsgi_app(environ, start_response) ❷

```

- ❶ This is line number 1973, identified in the debugger.
- ❷ This is line number 1985, also identified in the debugger. The WSGI server would receive the Flask instance as an application, and call it once for every request—by using the debugger, we’ve found the entry point for the code.

We’re using the debugger in the same way as we used the call graph with HowDoI—by following function calls—which is also the same thing as reading through code directly. The value of using the debugger is that we avoid looking at all of the additional code that may distract or confuse us. Use the approach that is most effective for you.

After going up the stack using `u`, we can go back down the stack using `d` and will end up back at the breakpoint, labeled with the `*** Newest frame`:

```

> /[ ... truncated path ...]/flask/examples/flaskr/flaskr.py(74)show_entries()
-> db = get_db()
(Pdb) d
*** Newest frame

```

From there, we can advance through a function call with the `n` (*next*) command, or advance in as short a step as possible with the `s` (*step*) command:

```

(Pdb) s
--Call--
> /[ ... truncated path ...]/flask/examples/flaskr/flaskr.py(55)get_db()
-> def get_db():
(Pdb) s
> /[ ... truncated path ...]/flask/examples/flaskr/flaskr.py(59)get_db()
-> if not hasattr(g, 'sqlite_db'): ❶
###~
###~ ... do a dozen steps to create and return the database connection...

```

```

##~~~
-> return g.sqlite_db
(Pdb) n
> /[ ... truncated path ... ]/flask/examples/flaskr/flaskr.py(75)show_entries()
-> cur = db.execute('select title, text from entries order by id desc')
(Pdb) n
> /[ ... truncated path ... ]/flask/examples/flaskr/flaskr.py(76)show_entries()
-> entries = cur.fetchall()
(Pdb) n
> /[ ... truncated path ... ]/flask/examples/flaskr/flaskr.py(77)show_entries()
-> return render_template('show_entries.html', entries=entries) ❷
(Pdb) n
--Return--

```

There's a lot more, but it's tedious to show. What we get out of it is:

- ❶ Awareness of the `Flask.g` object. A little more digging reveals it is the *global* (actually local to the Flask instance) context. It exists to contain database connections and other persistent things like cookies that need to survive outside of the life of the methods in the Flask class. Using a dictionary like this keeps variables out of the Flask app's namespace, avoiding possible name collisions.
- ❷ The `render_template()` function isn't much of a surprise, but it's at the end of the function definition in the *flaskr.py* module, meaning we're essentially done—the return value goes back to the calling function from the Flask instance that we saw when traversing up the stack. So we're skipping the rest.

The debugger is useful local to the place that you're inspecting, to find out precisely what's happening before and after an instant, the user-selected breakpoint, in the code. One of the big features is the ability to change variables on the fly (any Python code works in the debugger) and then continue on stepping through the code.

Logging in Flask

Diamond has an example of logging in an application, and Flask provides one of logging in a library. If all you want to do is avoid “no handler found” warnings, search for “logging” in the Requests library (*requests/requests/__init__.py*). But if you want to provide some logging support within your library or framework, Flask’s logging provides a good example to follow.

Flask-specific logging is implemented in *flask/flask/logging.py*. It defines the logging format strings for production (with logging level ERROR) and for debugging (with logging level DEBUG), and follows the advice from the [Twelve-Factor App](#) to log to streams (which direct to one of *wsgi.errors* or *sys.stderr*, depending on the context).

The logger is added to the main Flask application in *flask/flask/app.py* (the code snippet skips over anything that’s not relevant in the file):

```
# a lock used for logger initialization
_logger_lock = Lock() ❶

class Flask(_PackageBoundObject):

    ##~~ ... skip other definitions

    #: The name of the logger to use. By default the logger name is the
    #: package name passed to the constructor.
    #:
    #: .. versionadded:: 0.4
    logger_name = ConfigAttribute('LOGGER_NAME') ❷

    def __init__(self, import_name, static_path=None, static_url_path=None,
        ##~~ ... skip the other arguments ...
        root_path=None):
        ##~~ ... skip the rest of the initialization
        # Prepare the deferred setup of the logger.
        self._logger = None ❸
        self.logger_name = self.import_name

    @property
    def logger(self):
        """A :class:`logging.Logger` object for this application. The
        default configuration is to log to stderr if the application is
        in debug mode. This logger can be used to (surprise) log messages.
        Here some examples::

            app.logger.debug('A value for debugging')
            app.logger.warning('A warning occurred (%d apples)', 42)
```

```

        app.logger.error('An error occurred')

... versionadded:: 0.3
"""
if self._logger and self._logger.name == self.logger_name:
    return self._logger ❹
with _logger_lock: ❺
    if self._logger and self._logger.name == self.logger_name:
        return self._logger
    from flask.logging import create_logger
    self._logger = rv = create_logger(self)
    return rv

```

- ❶ This lock is used toward the end of the code. Locks are objects that can only be possessed by one thread at a time. When it is being used, any other threads that want it must block.
- ❷ Like Diamond, Flask uses the configuration file (with sane defaults, that aren't shown here, so the user can simply do nothing and get a reasonable answer) to set the logger name.
- ❸ The Flask application's logger is initially set to none so that it can be created later (in step ❺).
- ❹ If the logger exists already, return it. The property decoration, like earlier in this chapter, exists to prevent the user from inadvertently modifying the logger.
- ❺ If the logger doesn't exist yet (it was initialized to None), then use the lock created in step ❶ and create it.

Style Examples from Flask

Most of the style examples from [Chapter 4](#) have already been covered, so we'll only discuss one style example for Flask—the implementation of Flask's elegant and simple routing decorators.

Flask's routing decorators (beautiful is better than ugly)

The routing decorators in Flask add URL routing to target functions, like this:

```

@app.route('/')
def index():
    pass

```

The Flask application will, when dispatching a request, use URL routing to identify the correct function to generate the response. The decorator syntax keeps the routing code logic out of the target function, keeps the function flat, and is intuitive to use.

It's also not necessary—it exists only to provide this API feature to the user. Here is the source code, a method in the main Flask class defined in `flask/flask/app.py`:

```
class Flask(_PackageBoundObject): ❶
    """The flask object implements a WSGI application ...
    ... skip everything else in the docstring ...
    """
    """~ ... skip all but the routing() method.

    def route(self, rule, **options):
        """A decorator that is used to register a view function for a
        given URL rule. This does the same thing as :meth:`add_url_rule`
        but is intended for decorator usage::

            @app.route('/')
            def index():
                return 'Hello World'

        ... skip the rest of the docstring ...
        """
        def decorator(f): ❷
            endpoint = options.pop('endpoint', None)
            self.add_url_rule(rule, endpoint, f, **options) ❸
            return f
        return decorator
```

- ❶ The `_PackageBoundObject` sets up the file structure to import the HTML templates, static files, and other content based on configuration values specifying their location relative to the location of the application module (e.g., `app.py`).
- ❷ Why not name it `decorator`? That's what it does.
- ❸ This is the actual function that adds the URL to the map containing all of the rules. The only purpose of `Flask.route` is to provide a convenient decorator for library users.

Structure Examples from Flask

The theme for both of the structure examples from Flask is modularity. Flask is intentionally structured to make it easy to extend and modify almost everything—from the way that JSON strings are encoded and decoded (Flask supplements the standard library's JSON capability with encodings for datetime and UUID objects) to the classes used when routing URLs.

Application specific defaults (simple is better than complex)

Flask and Werkzeug both have a `wrappers.py` module. The reason is to add appropriate defaults for Flask, a framework for web applications, on top of Werkzeug's more

general utility library for WSGI applications. Flask subclasses Werkzeug's Request and Response objects to add specific features related to web applications. For example, the Response object in `flask/flask/wrappers.py` looks like this:

```
from werkzeug.wrappers import Request as RequestBase, Response as ResponseBase
##~~ ... skip everything else ...

class Response(ResponseBase): ❶
    """The response object that is used by default in Flask. Works like the
    response object from Werkzeug but is set to have an HTML mimetype by
    default. Quite often you don't have to create this object yourself because
    :meth:`~flask.Flask.make_response` will take care of that for you.

    If you want to replace the response object used you can subclass this and
    set :attr:`~flask.Flask.response_class` to your subclass. ❷
    """
    default_mimetype = 'text/html' ❸
```

- ❶ Werkzeug's Response class is imported as ResponseBase, a nice style detail that makes its role obvious and allows the new Response subclass to take its name.
- ❷ The ability to subclass `flask.wrappers.Response`, and how to do it, is documented prominently in the docstring. When features like this are implemented, it's important to remember the documentation, or users won't know the possibility exists.
- ❸ This is it—the only change in the Response class. The Request class has more changes, which we're not showing to keep the length of this chapter down.

This small interactive session shows what changed between Flask's and Werkzeug's Response classes:

```
>>> import werkzeug
>>> import flask
>>>
>>> werkzeug.wrappers.Response.default_mimetype
'text/plain'
>>> flask.wrappers.Response.default_mimetype
'text/html'
>>> r1 = werkzeug.wrappers.Response('hello', mimetype='text/html')
>>> r1.mimetype
u'text/html'
>>> r1.default_mimetype
'text/plain'
>>> r1 = werkzeug.wrappers.Response('hello')
>>> r1.mimetype
'text/plain'
```

The point of changing the default mimetype was just to make a little less typing for Flask users when building response objects that contain HTML (the expected use with Flask). Sane defaults make your code much, much easier for the average user.

Sane Defaults Can Be Important

Sometimes defaults matter a lot more than just for ease of use. For example, Flask sets the key for sessionization and secure communication to Null by default. When the key is null, an error will be raised if the app attempts to start a secure session. Forcing this error means users will make their own secret keys—the other (bad) options would be to either silently allow a null session key and insecure sessionization, or to provide a default key like *mysecretkey* that would invariably not be updated (and thus be used in deployment) by many.

Modularity (also one honking great idea)

The docstring for `flask.wrappers.Response` let users know that they could subclass the Response object and use their newly defined class in the main Flask object.

In this excerpt from `flask/flask/app.py`, we highlight some of the other modularity built into Flask:

```
class Flask(_PackageBoundObject):
    """ ... skip the docstring ...
    """

    #: The class that is used for request objects. See :class:`~flask.Request`
    #: for more information.
    request_class = Request ❶

    #: The class that is used for response objects. See
    #: :class:`~flask.Response` for more information.
    response_class = Response ❷

    #: The class that is used for the Jinja environment.
    #:
    #: .. versionadded:: 0.11
    jinja_environment = Environment ❸

    ##~ ... skip some other definitions ...

    url_rule_class = Rule
    test_client_class = None
    session_interface = SecureCookieSessionInterface()

    ##~ .. etc. .. ❹
```

❶ Here's where the custom Request class can be substituted.

- ❷ And here is the place to identify the custom `Response` class. These are class attributes (rather than instance attributes) of the `Flask` class, and are named in a clear way that makes it obvious what their purpose is.
- ❸ The `Environment` class is a subclass of Jinja2's `Environment` that has the ability to understand Flask Blueprints, which make it possible to build larger, multifile Flask applications.
- ❹ There are other modular options that are not shown because it was getting repetitive.

If you search through the `Flask` class definition, you can find where these classes are instantiated and used. The point of showing it to you is that these class definitions didn't have to be exposed to the user—this was an explicit structural choice done to give the library user more control over how `Flask` behaves.

When people talk about `Flask`'s modularity, they're not just talking about how you can use any database backend you want, they're also talking about this capability to plug in and use different classes.

You've now seen some examples of well-written, very Zen Python code.

We highly recommend you take a look at the full code of each of the programs discussed here: the best way to become a good coder is to read great code. And remember, whenever coding gets tough, *use the source, Luke!*